
On Polynomial Methods in Combinatorics

Conrad Crowley
118316041

Supervisor: Dr. Marco Vitturi
Second Reader: Dr. Andrei Mustata

Final Year Project 2022



Abstract

Polynomial methods are an emerging set of techniques in Combinatorics which use elementary facts about polynomials to control the size of collections of objects with a certain structure. We present an exposition on polynomial methods via a collection of results that can be elegantly proven with them, ranging from the solution to the Finite Field Kakeya Problem to bounds on the number of circle tangencies in the plane. For the latter problem we also provide a new proof.

Acknowledgements

I am greatly indebted to my supervisor Marco Vitturi for his time, guidance, and enthusiasm throughout this research endeavour. He has been extremely helpful throughout, far exceeding any expectations that could be reasonably set of a Bachelor's thesis supervisor.

I would like to thank Anca Mustata for some helpful conversations regarding the Real Algebraic Geometry considerations of Chapter 6.

I would also like to thank the second reader Andrei Mustata for taking the time to read and evaluate this (rather long) manuscript.

Further, gratitude must be expressed for the UCC Mathematics department without whom this extremely enjoyable experience would not have been possible. Special thanks should be given to the module coordinator Kevin Hayes who has been working diligently behind the scenes throughout the year to ensure this module progressed smoothly.

Table of Contents

Table of Contents	ii
List of Figures	iii
1 Introduction	1
1.1 Why Polynomials?	1
1.2 Notation	2
2 The Kakeya Problem in Finite Fields	3
2.1 Combinatorial Attempts	5
2.1.1 Bush Argument	6
2.1.2 Hairbrush Argument	7
2.2 Dvir's Proof	8
3 Cauchy-Davenport Theorem	12
3.1 Combinatorial Nullstellensatz	12
3.2 Proof of Cauchy-Davenport Theorem	13
4 The Joints Problem	15
4.1 Examples	15
4.2 Solution of the Joints Problem	16
5 Szemerédi-Trotter Theorem	18
5.1 The Trivial Bounds	18
5.2 Examples	20
5.3 Ham Sandwich Theorems	22
5.4 Proof of the Szemerédi-Trotter Theorem	26
6 Counting Circle Tangencies	29
6.1 Lifting of circles to \mathbb{R}^3	31
6.2 Ellenberg-Solymosi-Zahl's Proof	31
6.3 New Proof via Polynomial Partitioning	34
6.4 The Case of Sphere Tangencies in \mathbb{R}^3	39

References	43
Appendices	44
A Proof of the Borsuk-Ulam Theorem	44

List of Figures

2.1	An example of a Kakeya set (shaded) in \mathbb{R}_3^2	5
4.1	A $N \times N$ layer of our grid.	16
5.1	Example 5.3 in the case $N = 3$. (with the x, y axis reversed)	21
6.1	A collection of circles in \mathbb{R}^2 with N^2 tangencies.	30
6.2	Two collections of N spheres in \mathbb{R}^3 , whose union satisfies the non-degeneracy condition the earlier section yet achieves N^2 tangencies.	40

Chapter 1

Introduction

The following is a short exposition of polynomial methods in Combinatorics. Polynomial methods are a collection of techniques which use polynomial interpolation and rigidity properties of polynomials to control the size of collections of objects with a certain structure. The first example of this technique was presented in the 1990s in [Alo99], which we examine in detail in Chapter 3. The modern conception of the polynomial method was pioneered by Dvir in 2008 (See [Dvi09]), where he produced a remarkably short resolution to the finite field analogue of the Kakeya Conjecture which provided a new framework and enthusiasm for the polynomial method in combinatorial problems. We will explore this proof in the next chapter.

The most striking feature of the following proofs is that they leverage certain properties of polynomials in problems which on the surface appear not to have anything to do with polynomials. Generally, extremal configurations of these problems tend to admit a lot of algebraic structure and this is exactly what these methods exploit using polynomials.

1.1 Why Polynomials?

It is perhaps wise to discuss here what features of polynomials make them particularly powerful when dealing with problems in Combinatorics. Polynomials are perhaps some of the simplest of mathematical objects, as once we define a field they are simply a combination of the addition and multiplication operation between elements. It is not immediately obvious why such simple objects may prove to be so useful.

There are two key properties of polynomials that this collection of methods exploit. Firstly, we use the fact that there are roughly $\sim D^n$ coefficients of a polynomial in n variables of degree at most D (See Lemma 2.2.2). This is utilised in an essential manner when we try to find polynomials that contain objects in their zero set. We can contain a set of size M in the zero set of a polynomial with degree at most $O(M^{1/n})$. In other words, we have a lot of flexibility in choosing a polynomial. Secondly, and in sharp contrast, polynomials behave extremely rigidly when restricted to lines. We mean by this that the zero set of a polynomial of degree D can intersect a line in at most D points if the line is not

contained within said zero set. The gap between this flexibility of choosing a polynomial and rigidity of restricting to lines provides us with a surprisingly powerful technique.

Another striking thing about the method is the non-constructive manner in which the polynomials are usually used. We often cannot explicitly find a satisfactory polynomial to use for our purposes, instead we opt to use arguments from Linear Algebra to establish the existence of a polynomial with such properties. This is reminiscent of other methods in Combinatorics, such as the Probabilistic Method or the Topological Method (See [Alo03] for a survey of these methods).

1.2 Notation

We introduce some convenient notation here. We write that $A \lesssim_n B$ to mean that there exists some constant $C(n)$ which depends on n such that $A \leq C(n)B$. Further, we write that $A \sim_n B$ if $A \lesssim_n B$ and $B \lesssim_n A$.

We write $\text{Poly}_D(\mathbb{K}^n)$ to represent the space of polynomials in n variables with coefficients in a field \mathbb{K} and degree at most D .

The indicator function $\mathbb{1}$ is defined on logical statements X as follows:

$$\mathbb{1}[X] = \begin{cases} 1 & \text{if } X \text{ is true,} \\ 0 & \text{if } X \text{ is false.} \end{cases}$$

For any function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ let us denote the zero set of f by $Z(f) = \{x \in \mathbb{R}^n \mid f(x) = 0\}$. We borrow from Computer Science the big O notation. For functions $f, g : \mathbb{N}^+ \rightarrow \mathbb{R}$ we write:

$$\begin{aligned} f(N) = O(g(N)) &\iff \exists N_0, M \in \mathbb{N} \text{ such that } f(n) \leq Mg(n) \forall n > N_0 \\ f(N) = \Omega(g(n)) &\iff g(N) = O(f(N)). \end{aligned}$$

These can be thought of as asymptotic upper and lower bounds respectively.

Chapter 2

The Kakeya Problem in Finite Fields

Before we can discuss the Kakeya problem in finite fields, and its rather surprising resolution, we ought to first discuss the origin and history of the problem. Work on the Kakeya problem can be traced back to the Russian mathematician Abram Besicovitch in 1917. While working on a problem in Riemann integration, Besicovitch was led to consider the question of the existence of planar sets of measure zero which contain a line segment in every direction. In 1920, Besicovitch constructed such a set and published in a Russian Journal.

However, 1917 was a turbulent year as it marked the end of the Russian Empire and the start of the Russian civil war. Due to this and the ensuing blockade of Russian ports there was scarce communication with the outside world. Thus Besicovitch could not have known of a Japanese mathematician Kakeya who asked also in 1917 a related question: What is the smallest area of a convex set within which one can rotate a needle by 180 degrees in the plane? Julius Pal answered this question in 1921 with the equilateral triangle in [Pal20]. The more interesting problem obtained by dropping the convexity condition remained open. In 1924, after leaving the newly formed Soviet Union for Copenhagen, Besicovitch discovered this problem and by modifying his previous construction produced a solution in 1925. This led to the more general questions being asked about Kakeya sets in higher dimensions.

Definition 2.0.1 (Kakeya Set in \mathbb{R}^n). A **Kakeya set** is a set $A \subset \mathbb{R}^n$ that contains a unit segment in every direction.

Besicovitch's construction showed that these sets can have arbitrarily small measures, even attaining zero, in \mathbb{R}^2 . Further, a straightforward process allows to extend the construction of these measure-zero sets to the case of arbitrary dimension.

Given that Kakeya sets can have measure zero, the natural question then arises, are Kakeya sets fractal in nature? This leads one to consider the dimension of Kakeya sets as a quantification of their fractal nature. There are many notions of dimensions that can be

investigated, but we restrict ourselves to the Minkowski and Hausdorff dimensions.

Definition 2.0.2 (Minkowski Dimension). Given a set $S \subset \mathbb{R}^n$, define $N(\varepsilon)$ to be the number of cubes of side length ε required to cover the set. The **Minkowski Dimension** of the set S is then defined as

$$\dim_M(S) = \lim_{\varepsilon \rightarrow 0} \frac{\log(N(\varepsilon))}{\log(1/\varepsilon)}.$$

If this limit does not exist, one can still define the upper and lower Minkowski dimensions, $\dim_{M_{\text{upper}}}$ and $\dim_{M_{\text{lower}}}$, by taking the limit superior and limit inferior respectively.

Definition 2.0.3 (Hausdorff Dimension). We define the d -dimensional Hausdorff measure of a set $S \subset \mathbb{R}^n$ as

$$\mathcal{H}^d(S) = \liminf_{r \rightarrow 0} \left\{ \sum_i r_i^d : \text{there is a countable cover of } S \text{ by balls with radii } 0 < r_i < r \right\}.$$

Then we can define the **Hausdorff dimension** of the set S to be

$$\dim_H(S) = \inf\{d \geq 0 : \mathcal{H}^d(S) = 0\}.$$

These dimensions are related by the following inequality when they are all defined:

$$\dim_H \leq \dim_{M_{\text{lower}}} \leq \dim_{M_{\text{upper}}}.$$

In 1971, Davies published [Dav71] showing that although the measure of a Kakeya set in \mathbb{R}^2 can be arbitrarily small, it must have Hausdorff and Minkowski dimension of 2. This resulted in the following conjectures:

Conjecture 1 (Kakeya Conjecture for the Minkowski Dimension). Let A be a Kakeya set in \mathbb{R}^n . Then $\dim_M(A) = n$.

Conjecture 2 (Kakeya Conjecture for the Hausdorff Dimension). Let A be a Kakeya set in \mathbb{R}^n . Then $\dim_H(A) = n$.

In a survey on the problem, (see [Wol99]) Wolff proposed a finite analogue to the Kakeya Conjecture. We begin with some preliminaries so we can formally state this analogue.

Definition 2.0.4 (Finite Field). A **finite field** \mathbb{F}_q is a field of finite cardinality. The cardinality $|\mathbb{F}_q| = q$ of a finite field is called the **order** of the finite field.

It is known that finite field of order q exists if and only if $q = p^k$ for some prime p and integer k . For the rest of this chapter we let q be of this form.

Lemma 2.0.1. *Each element X in a finite field \mathbb{F} satisfies the identity:*

$$X^{|\mathbb{F}|} - X = 0$$

identically in \mathbb{F} .

Proof. This lemma follows immediately from Fermat's Little Theorem. \square

A Kakeya set in \mathbb{F}_q^n is a set that contains a line in every direction. In analogy to the Euclidean case, we formally define lines in \mathbb{F}_q^n as the sets of the form

$$\ell_{x,y} = \{x + ty : t \in \mathbb{F}_q\}$$

for some fixed $x, y \in \mathbb{F}_q^n$ with $y \neq 0$. It should be noted that a line in a finite field \mathbb{F}_q^n contains exactly $|\mathbb{F}_q|$ points. Formally, directions in \mathbb{F}_q^n can be identified using the projective space $\mathbb{P}\mathbb{F}_q^n = \mathbb{F}_q^n / \mathbb{F}_q^\times$. In this projective space, two lines y and y' are equivalent if y' is a translation of y . The finite analogue to Conjectures 1 and 2 is:

Conjecture 3 (Kakeya Conjecture in Finite Fields). If $A \subset \mathbb{F}_q^n$ contains a line in every direction, then $|A| \gtrsim_n |\mathbb{F}_q|^n$.

This conjecture had a significant influence on the subject, inspiring work on the sum-product phenomenon in finite fields. From its postulation in 1999, little progress was made in the following years and it was assumed that the problem was roughly as difficult as the Euclidean case.

In 2008, Dvir published a remarkably simple proof of Conjecture 3 using elementary facts about polynomials (See [Dvi09]). This proof revitalised interest in the polynomial method, and we shall explore its details later in this chapter.




Figure 2.1: An example of a Kakeya set (shaded) in \mathbb{F}_3^2 .

2.1 Combinatorial Attempts

To truly show the power of the polynomial method employed by Dvir, we shall first explore purely combinatorial attempts at estimating the sizes of Kakeya sets. All the following estimates are much weaker than the bound in Theorem 3. We fix a finite field $\mathbb{F} = \mathbb{F}_{p^k}$ where p is a prime.

Our first estimate shows that Conjecture 3 is true for the $n = 2$ case.

Lemma 2.1.1. *Suppose $s \leq |\mathbb{F}|$. If l_1, \dots, l_s are distinct lines in \mathbb{F}^n , then their union has cardinality at least $(1/2)|\mathbb{F}|s$.*

In particular, if $A \subset \mathbb{F}^n$ is a Kakeya set, then we have

$$|A| \gtrsim |\mathbb{F}|^2.$$

Proof. We add the lines together one at a time and track the cardinality of their union. The first line contains $|\mathbb{F}|$ points. The second line must contain at least $|\mathbb{F}| - 1$ points not in the first line. Similarly the third line must contain at least $|\mathbb{F}| - 2$ points not in the first two lines, and so on. Thus the number of distinct points in the union of all s lines is given by

$$\sum_{i=1}^s |\mathbb{F}| - i + 1 > (1/2)|\mathbb{F}|s.$$

A Kakeya set always contains at least $|\mathbb{F}|$ lines, so setting $s = |\mathbb{F}|$ above yields

$$|A| > (1/2)|\mathbb{F}|^2.$$

□

2.1.1 Bush Argument

Bourgain produced one of the first non-trivial estimates of the dimension in his work [Bou91]. We present the finite field analogue of his argument here (See [Gut16]). This argument is known as the “bush” argument as we consider a high-multiplicity point through which pass many lines, forming a “bush” around that point.

Theorem 2.1.2 (Bush Argument). *If $A \subset \mathbb{F}^n$ is a Kakeya set, then we have*

$$|A| \gtrsim |\mathbb{F}|^{\frac{n+1}{2}}.$$

Proof. Let μ be a fixed multiplicity parameter to be chosen later. Either there exists a point $p \in A$ such that there are μ lines passing through p , or else every point in A has less than μ lines passing through it.

In the first case, since the lines have distinct directions they must become disjoint when p is removed. Hence

$$|A| \gtrsim \mu|\mathbb{F}|.$$

In the latter case, by double counting we have

$$\begin{aligned} \sum_{\ell \subset A} \sum_{p \in A} \mathbb{1}[p \in \ell] &= \sum_{\ell \subset A} |\mathbb{F}| = |\mathbb{F}|^{n-1} |\mathbb{F}| \\ \sum_{p \in A} \sum_{\ell \subset A} \mathbb{1}[p \in \ell] &\leq \sum_{p \in A} \mu = |A| \mu \\ \implies \frac{|\mathbb{F}|^n}{\mu} &\leq |A|. \end{aligned}$$

Now we optimise the two lower bounds by choosing $\mu \sim |\mathbb{F}|^{\frac{n-1}{2}}$, so that in either of the above cases we obtain

$$|A| \gtrsim |\mathbb{F}| |\mathbb{F}|^{\frac{n-1}{2}} \sim |\mathbb{F}|^{\frac{n+1}{2}}.$$

□

2.1.2 Hairbrush Argument

In [Wol95], Wolff presented the “hairbrush” argument where he made an incremental improvement over the argument of Bourgain. It is similar in spirit to the “bush” argument above, however instead of using just one point of high-multiplicity we instead consider a line (or “stem”) containing many points of high-multiplicity.

Theorem 2.1.3 (Hair Brush Argument). *If $A \subset \mathbb{F}^n$ is a Kakeya set, then we have*

$$|A| \gtrsim |\mathbb{F}|^{\frac{n+2}{2}}.$$

Proof. Let μ be a fixed multiplicity parameter to be chosen later. We say a line ℓ is μ -rich if for at least $|\mathbb{F}|/2$ points $p \in \ell$ there are μ lines distinct from ℓ in A passing through p . Either there exists a μ -rich line or there does not.

Suppose a μ -rich line exists, and denote this line by ℓ_μ . Consider the family Π of 2-dimensional planes passing through ℓ_μ . If a line ℓ intersects ℓ_μ then there exists a unique plane $\pi \in \Pi$ such that $\ell, \ell_\mu \in \pi$. Let $\mathcal{L}_\pi := \{\ell \subset \pi \mid \ell \cap \ell_\mu \neq \emptyset\}$. The set $A \cap \pi$ is a set in (an isomorphic copy of) \mathbb{F}^2 that contains at least $|\mathcal{L}_\pi|$ lines, and hence by Lemma 2.1.1 we have

$$|A \cap \pi| \gtrsim |\mathcal{L}_\pi| |\mathbb{F}|.$$

Thus,

$$|A| \geq \sum_{\pi \in \Pi} |(A \cap \pi) \setminus \ell_\mu| \gtrsim |\mathbb{F}| \sum_{\pi \in \Pi} |\mathcal{L}_\pi| \gtrsim \mu |\mathbb{F}|^2,$$

where the last inequality comes from the fact ℓ_μ intersects at least $\mu |\mathbb{F}|/2$ lines, so since the planes Π foliate \mathbb{F}^n , the union $\cup_{\pi \in \Pi} \mathcal{L}_\pi$ contains at least all of these lines.

Suppose there does not exist a μ -rich line. Let

$$A' = \{p \in A \mid p \text{ belongs to strictly less than } \mu \text{ lines in } A\}.$$

Since we assume there does not exist a μ -rich line in A , for any line $\ell \subset A$ we have $|A' \cap \ell| > |F|/2$. Proceeding by double counting,

$$\begin{aligned} |A| &\geq |A'| \geq \frac{1}{\mu} \sum_{p \in A'} \sum_{\ell \subset A} \mathbb{1}[p \in \ell] \\ &= \frac{1}{\mu} \sum_{\ell \subset A} |A' \cap \ell| \gtrsim \frac{1}{\mu} |\mathbb{F}|^{n-1} |\mathbb{F}| \\ \implies |A| &\gtrsim \frac{|\mathbb{F}|^n}{\mu}. \end{aligned}$$

Now optimising the two lower bounds by choosing $\mu \sim |\mathbb{F}|^{\frac{n-2}{2}}$, so that in either of the above cases we obtain the bound

$$|A| \gtrsim |\mathbb{F}|^2 |\mathbb{F}|^{\frac{n-2}{2}} \sim |\mathbb{F}|^{\frac{n+2}{2}}.$$

□

2.2 Dvir's Proof

Again we fix $\mathbb{F} = \mathbb{F}_{p^q}$ for some prime p .

Theorem 2.2.1 (Kakeya Conjecture in Finite Fields). *If $A \subset \mathbb{F}^n$ contains a line in every direction, then $|A| \geq \frac{1}{n!} |\mathbb{F}|^n$.*

We shall prove this theorem via three surprisingly elementary lemmas. The general strategy of the proof is to assume the size of a Kakeya set A is small and find a low-degree non-zero polynomial that interpolates (contains in its zero set) the points in A . We then show that due to the structure of the Kakeya set the polynomial's zero set must contain many more points. This contradicts the fact our polynomial is non-zero and of low-degree which establishes the Kakeya Conjecture over finite fields.

The first lemma is a fundamental one to the polynomial method, exploiting the flexible interpolation properties of polynomials. Consider the problem of finding a non-zero polynomial in $\mathbb{R}[X, Y, Z]$ of the minimal possible degree that vanishes at the points $(j, k, 2^{j^k})$ for $1 \leq i \leq 10^6$. If we try to find an explicit polynomial we may produce

$$P(X, Y, Z) = \prod_{j,k=1}^{10^6} (X - j)(Y - k)$$

or perhaps

$$Q(X, Y, Z) = \prod_{j,k=1}^{10^6} (Z - 2^{j^k}).$$

Both P and Q vanish on our points and it is difficult to explicitly produce a polynomial of much lower degree that also satisfies this property. One might consider taking the product of linear factors in $\mathbb{R}[X, Y, Z]$, constructing the polynomial to vanish on sets of planes defined by three points in the collection. Such a polynomial has a degree of $\lceil \frac{1}{3}10^{12} \rceil$, which is slight improvement on Q which has a degree of 10^{12} . Naively one might conjecture that the degree of any polynomial with this property is of the order of 10^{12} . Remarkably, the following lemma shows that we can find such a polynomial with a degree of about ~ 18000 . This polynomial has roughly $\sim 18000^3$ coefficients and as such the polynomial is quite difficult to explicitly write, but we can establish the existence of such a polynomial non-constructively.

Lemma 2.2.2 (Parameter Counting). *Let \mathbb{K} be a (not necessarily finite) field. If $A \subset \mathbb{K}^n$ and $|A| < \binom{n+D}{n}$, there exists a non-zero polynomial $P(X_1, \dots, X_n)$ of degree at most D that vanishes on A .*

Proof. We first show the dimension of $\text{Poly}_D(\mathbb{K}^n)$ is $\binom{D+n}{n}$. A basis for $\text{Poly}_D(\mathbb{K}^n)$ is given by monomials of the form $X_1^{D_1} \dots X_n^{D_n}$, where $\sum_i D_i \leq D$, hence we just need to count the number of monomials.

We can map a monomial $X_1^{D_1} \dots X_n^{D_n}$ to a string of D \star 's and n $|$'s as follows. Begin with D_1 \star 's, then place one $|$. We put now D_2 \star 's, and place a second $|$. We continue until we have placed D_n \star 's followed by an n^{th} $|$. Finally we place $D - \sum_i D_i$ \star 's. This is a bijective map between the monomials in $\text{Poly}_D(\mathbb{K}^n)$ and all the strings of D \star 's and n $|$'s. To assemble such a string we distribute the n $|$'s over $n + D$ spaces, and then fill the remainder with \star 's. Hence we get the binomial coefficient

$$\text{Poly}_D(\mathbb{K}^n) = \binom{n+D}{n}.$$

Now let $p_1, \dots, p_{|A|}$ be the points of A . We consider the evaluation map $E : \text{Poly}_D(\mathbb{K}^n) \rightarrow \mathbb{K}^{|A|}$ defined by

$$E(Q) = (Q(p_1), \dots, Q(p_{|A|})).$$

This map is clearly linear. Its kernel $\ker E$ is exactly the set of polynomials in $\text{Poly}_D(\mathbb{K}^n)$ that vanish on A . By assumption, the dimension of $\text{Poly}_D(\mathbb{K}^n)$ is greater than $|A|$, so the dimension of the domain of E is greater than the codomain of E . By the rank-nullity theorem, we conclude E must have a non-trivial kernel. Thus there exists a non-zero polynomial $P \in \text{Poly}_D(\mathbb{K}^n)$ that vanishes on A . \square

Note that if $D = |\mathbb{F}| - 1$, and $|A| \leq \binom{|\mathbb{F}|+n-1}{|\mathbb{F}|-1} = \binom{|\mathbb{F}|+n-1}{n}$ we have a polynomial of degree at most $|\mathbb{F}| - 1$ that vanishes on A . Since $\frac{|\mathbb{F}|^n}{n!} < \binom{|\mathbb{F}|+n-1}{n}$, we can certainly find such a polynomial when $|A| \leq \frac{|\mathbb{F}|^n}{n!}$. The restriction of $D = |\mathbb{F}| - 1$ is somewhat necessary due to the fact that polynomials with a factor of $X^{|\mathbb{F}|} - X$ vanish identically by Lemma 2.0.1.

Lemma 2.2.3. *Suppose $A \subset \mathbb{F}^n$ contains a line in every direction, and suppose that there exists a non-zero polynomial P with degree $D < |\mathbb{F}|$ that vanishes on A . Then there exists a non-zero degree D polynomial \bar{P} that vanishes everywhere on \mathbb{F}^n .*

Proof. Choose a line in A , say $\ell = \{x + ty : t \in \mathbb{F}\}$ with $x, y \in \mathbb{F}^n$ and $y \neq 0$. Now we consider the restriction of our polynomial P to the line ℓ , $P|_\ell$. Recall P is a sum of monomials, and we use multi-index notation here with $\alpha = (\alpha_1, \alpha_2, \dots, \alpha_n)$, $\alpha_i \in \mathbb{N} \cup \{0\}$ and $|\alpha| = \sum \alpha_i$. P can be written as

$$P(X_1, X_2, \dots, X_n) = \sum_{|\alpha| \leq D} c_\alpha X_1^{\alpha_1} X_2^{\alpha_2} \dots X_n^{\alpha_n}.$$

Now $P|_\ell$ can be written

$$P|_\ell = P(x + ty) = Q_{x,y}(t) = \sum_{|\alpha| \leq D} c_\alpha \prod_i (x_i + ty_i)^{\alpha_i}.$$

We now wish to examine the degree D term of Q , which consist in the terms of the expansion of the expression above obtained by the product of the ty_i terms when $|\alpha| = D$.

This is related to the degree D component Q_D of Q by the identity

$$Q_{x,y,D} = t^D Q_D(y) = t^D \sum_{|\alpha|=D} c_\alpha \prod_i y_i^{\alpha_i}.$$

Now if $P|_\ell$ vanishes everywhere on ℓ , since its dependence on t is given by a polynomial of degree less than $|\mathbb{F}|$, all its coefficients must be zero. In particular, we must have $Q_{x,y,D} = 0$. This is clear from the factor theorem, as we could write the roots of $P|_\ell$ as $(t - k_1)(t - k_2) \dots (t - k_{|\mathbb{F}|})$; but this contradicts the fact P is of degree $D < |\mathbb{F}|$.

Notice that $Q_{x,y,D}$ no longer depends on x , but on y alone. In particular $Q_D(y) = 0$, but y was an arbitrary non-zero element of \mathbb{F}^n , and $Q_D(y)$ also vanishes at zero, so it vanishes everywhere. Thus we can pick $\bar{P} = Q_D$, and we are done. \square

Finally, we show that any non-zero polynomial in $\mathbb{F}[X_1, \dots, X_n]$ with degree less than $|\mathbb{F}|$ cannot be zero everywhere in \mathbb{F}^n .

Lemma 2.2.4. *Let P be a non-zero polynomial on \mathbb{F}^n with degree less than $|\mathbb{F}|$. Then P does not vanish everywhere.*

Proof. We proceed by induction on n . For $n = 1$, a non-zero polynomial that vanishes everywhere has $|\mathbb{F}|$ roots, so must be at least of degree $|\mathbb{F}|$. Let us assume that the statement holds in \mathbb{F}^{n-1} , we now prove it must also hold for \mathbb{F}^n .

We proceed by contradiction, assuming that P vanishes everywhere. We let X_1, \dots, X_n be coordinates on \mathbb{F}^n , and we write P in the form

$$P(X_1, \dots, X_n) = \sum_{j=n}^{|\mathbb{F}|-1} P_j(X_1, \dots, X_{n-1}) X_n^j.$$

Each P_j are polynomials in X_1, \dots, X_{n-1} of degree less than $|\mathbb{F}|$. Fix X_1, \dots, X_{n-1} , and let X_n vary. Now we have a polynomial in X_n of degree less than $|\mathbb{F}|$ that vanishes for all $X_n \in \mathbb{F}$. By the base case this must be the zero polynomial. So each $P_j(x_1, \dots, x_{n-1}) = 0$ for all j and for all $(x_1, \dots, x_{n-1}) \in \mathbb{F}^{n-1}$. Now by induction on n , each P_j is the zero polynomial. Then P is the zero polynomial as well. \square

We now combine these lemmas to establish the Kakeya conjecture over finite fields.

Proof of Theorem 2.2.1. Assume $A \subset \mathbb{F}^n$ is a Kakeya set, and that $|A| \leq \frac{|\mathbb{F}|^n}{n!}$. Then by Lemma 2.2.2 we can find a non-zero polynomial of degree at most $|\mathbb{F}| - 1$, say P , that vanishes on A . Now by Lemma 2.2.3 there exists a non-zero polynomial \bar{P} that vanishes everywhere on \mathbb{F}^n , and has degree less than $|\mathbb{F}|$. Finally, Lemma 2.2.4 says that such a \bar{P} is necessarily the zero polynomial, a contradiction. We conclude that $|A| > \frac{|\mathbb{F}|^n}{n!}$, or in other words

$$|A| \gtrsim_n |\mathbb{F}|^n.$$

\square

Remark. *Theorem 2.2.1 lends a good degree of support for Conjectures 1 and 2 as it essentially rules out purely algebraic counter-examples.*

Chapter 3

Cauchy-Davenport Theorem

In this chapter we shall discuss applications to the exciting field of additive combinatorics, where in 1990 Alon provided perhaps the first example of a polynomial method in action (see [Alo99]). This chapter is distinct from the other chapters in this manuscript as it deals with applications outside a geometric setting. It showcases a different polynomial method to the previous chapter, employing instead a combinatorial version of Hilbert's Nullstellensatz. In [Mic10], Michalek produced a short, elementary, and direct proof of the combinatorial Nullstellensatz which we present here.

3.1 Combinatorial Nullstellensatz

The German term *Nullstellensatz* means “theorem about zeros”, so it should come as no surprise that the theorem is precisely a statement about the general size and shape of sets of zeros of a polynomial depending on its highest degree terms.

The main idea behind the Nullstellensatz comes from the fact that a polynomial in one variable of degree D cannot have more than D roots. As a consequence of this if we have a set of $D + 1$ elements they cannot all be roots of the same degree D polynomial. In the univariate case this is somewhat elementary, but it has a very non-trivial extension to grids of points in arbitrary dimensions.

Theorem 3.1.1 (Combinatorial Nullstellensatz). *Let \mathbb{K} be a (not necessarily finite) field, and let $P(X_1, \dots, X_n) \in \mathbb{K}[X_1, \dots, X_n]$ be a polynomial in n variables with coefficients in \mathbb{K} . Suppose the coefficient of $X_1^{k_1} X_2^{k_2} \dots X_n^{k_n}$ is non-zero and further suppose $\deg P = \sum_{i=1}^n k_i$, where each k_i is a non-negative integer.*

Then for any subsets A_1, \dots, A_n of \mathbb{K} satisfying $|A_i| > k_i$ for each $1 \leq i \leq n$ there exist $a_1 \in A_1, \dots, a_n \in A_n$ such that $P(a_1, \dots, a_n) \neq 0$.

Proof. We proceed by induction on $\deg P = D$. When $D = 1$, P is simply a linear combination of n variables say $P(X_1, \dots, X_n) = c_1 X_1 + \dots + c_n X_n$. Without loss of generality assume X_1 has a non-zero coefficient and consider the sets $A_i = \{a_i, b_i\}$. Suppose P at

the point (a_1, a_2, \dots, a_n) is zero. We can then determine

$$c_1 = -\frac{c_2 a_2 + c_3 a_3 + \dots + c_n a_n}{a_1}.$$

Now evaluating at (b_1, a_2, \dots, a_n) we see that this is zero only when $a_1 = b_1$, so our theorem holds.

Now let us assume the theorem holds for $\deg P = D - 1$, and prove it for $\deg P = D$. Suppose that P satisfies the assumptions of the theorem but $P(x) = 0$ for every $x \in A_1 \times \dots \times A_n$. Without loss of generality $k_1 > 0$. Fixing $a \in A_1$ we can write

$$P = (X_1 - a)Q + R \tag{3.1}$$

by the usual long division of polynomials. The degree of R in X_1 must be strictly less than $\deg(X_1 - a)$, so R is independent of X_1 . Thus it follows that Q must have a monomial with non-zero coefficient of the form $X_1^{k_1-1} X_2^{k_2} \dots X_n^{k_n}$ and $\deg(Q) = D - 1$.

Take any $x \in \{a\} \times A_2 \times \dots \times A_n$ and evaluate (3.1). Since $P(x) = 0$ it follows that $R(x) = 0$. Hence R vanishes identically on the slice $\{a\} \times A_2 \times \dots \times A_n$. Since R is independent of X_1 it must also vanish identically on $A_1 \times A_2 \times \dots \times A_n$. Now take any $x \in A_1 \setminus \{a\} \times A_2 \times \dots \times A_n$ and evaluate (3.1). Since the $(X_1 - a)$ term is non-zero, $Q(x) = 0$. So Q vanishes on all $x \in A_1 \setminus \{a\} \times A_2 \times \dots \times A_n$, which contradicts the inductive hypothesis. \square

3.2 Proof of Cauchy-Davenport Theorem

To showcase the usefulness of the Combinatorial Nullstellensatz, we shall prove a classical result in Additive Combinatorics.

Theorem 3.2.1 (Cauchy-Davenport Theorem). *Let A, B be non-empty subsets of \mathbb{Z}_p for some prime p . Define their sumset $A + B$ as*

$$A + B = \{x \in \mathbb{Z}_p \mid x = a + b \text{ for some } a \in A, b \in B\}.$$

Then we have

$$|A + B| \geq \min\{p, |A| + |B| - 1\}.$$

Proof. Let us tackle the two cases separately. First, assume that $\min\{p, |A| + |B| - 1\} = p$. Then if $|A| + |B| > p$, A and B must intersect by the pigeonhole principle. For any $g \in \mathbb{Z}_p$ denote the set $\{g - x \mid x \in B\} \subset \mathbb{Z}_p$ as $g - B$. Since $|g - B| = |B|$, we have that $g - B$ and A must intersect as well. Thus there exists some $a \in A, b \in B$ such that

$$g - b = a \implies g = a + b.$$

Our choice of g was arbitrary, so it follows that $A + B = \mathbb{Z}_p$ and hence $|A + B| = p$.

Now assume that $\min\{p, |A| + |B| - 1\} = |A| + |B| - 1$. Then if the theorem is false we have $|A + B| \leq |A| + |B| - 2$, so there exists some $C \subset \mathbb{Z}_p$ such that $A + B \subset C$ and $|C| = |A| + |B| - 2$. Now let us define a polynomial $f(X, Y) \in \mathbb{Z}_p[X, Y]$ as

$$f(X, Y) = \prod_{c \in C} (X + Y - c).$$

Since $A + B \subset C$, $f(a, b) = 0$ for all $(a, b) \in A \times B$. Further, the degree of f is $\deg f = |C| = |A| + |B| - 2$. We can now appeal to the Combinatorial Nullstellensatz to yield a contradiction. Let $k_1 = |A| - 1$, and $k_2 = |B| - 1$. Now $\deg f = k_1 + k_2$, and the coefficient of $X^{k_1}Y^{k_2}$ is $\binom{|A|+|B|-2}{|A|-1}$, which is non-zero in \mathbb{Z}_p as the binomial coefficient is not divisible by p if all its factors are less than p . Applying Theorem 3.1.1 we see that there must exist some $(a, b) \in A \times B$ such that $f(a, b) \neq 0$, a contradiction. \square

The Cauchy-Davenport Theorem (Theorem 3.2.1) is indeed tight as illustrated by the following example.

Example 3.1. Fix $b \in \mathbb{Z}_p$ and consider the subsets of \mathbb{Z}_p given by

$$\begin{aligned} A &= \{tb \mid 1 \leq t \leq m\}, \\ B &= \{t'b \mid 1 \leq t' \leq n\}. \end{aligned}$$

Notice that $|A| = m$ and $|B| = n$. Their sumset is given by

$$A + B = \{(t + t')b \mid 1 \leq t \leq m, 1 \leq t' \leq n\}.$$

The $(t + t')$ term takes on every value between 2 and $m + n$ at least once. Hence

$$|A + B| = m + n - 1 = |A| + |B| - 1.$$

Chapter 4

The Joints Problem

Definition 4.0.1. Let \mathcal{L} be a set of distinct lines in \mathbb{R}^n . A **joint** of \mathcal{L} is a point which lies in three non-coplanar lines of \mathcal{L} .

The joints problem consists in obtaining a sharp upper bound on the maximal number of joints that can be formed from a configuration of L distinct lines. We denote this quantity $J(L)$. In other words $J(L)$ is the supremum over all configurations of lines in \mathbb{R}^n of the number of joints.


The joints problem was first posed in 1990 by Chazelle et al in [CEG⁺90]. They focused on the 3-dimensional case of the problem, establishing a lower bound of $\Omega(N^{3/2})$ and an upper bound estimate of $O(N^{7/4})$. This upper bound exponent has fallen gradually throughout the years, the best result prior to the application of the polynomial method being due to Sharir and Feldman in [FS05] where they established an upper bound of $O(N^{112/69})$. Their proof uses an array of tools from Combinatorial Geometry, including forbidden subgraphs in extremal graph theory, space decomposition techniques, and some results from the geometry of lines in space.

4.1 Examples

We shall begin by examining an example based on a grid, in order to gain better intuition about the problem and to formulate a conjecture.

Example 4.1. Consider an $N \times N \times N$ regular grid of integer coordinates in \mathbb{R}^3 . We shall give a collection of lines such that each point of this grid is a joint for the collection. Let \mathcal{L} be the collection of all lines parallel to any of the Cartesian axes that intersect a point in this grid. For each horizontal $N \times N$ layer, there are $N + N = 2N$ such lines that intersect our grid. There are N layers, so we obtain $2N^2$ distinct lines in this manner. Finally we need to account for the N^2 lines perpendicular to the horizontal $N \times N$ layers. This leaves us with $|\mathcal{L}| = 3N^2$ lines forming N^3 joints. The number of joints is thus $\sim |\mathcal{L}|^{3/2}$.

We can extend this example to higher dimensional grids easily.

Figure 4.1: A $N \times N$ layer of our grid.

Example 4.2. If we have an $\overbrace{N \times \cdots \times N}^{n \text{ Dimensions}}$ regular grid of integer coordinates in \mathbb{R}^n , we can construct an example by a straightforward extension of the above example. Each additional dimension increases the number of lines by a factor of N (this can be seen by considering each new dimension as a layering of the previous set along the new axis). Thus we can see that $\sim N^{n-1}$ lines form N^n joints in this manner. So the number of joints is $\sim |\mathcal{L}|^{\frac{n}{n-1}}$.

It turns out that the examples illustrated above provide asymptotically maximal configurations, that is, the best configurations disregarding the best constant C such that $J(L) \leq CL^{\frac{n}{n-1}}$.

4.2 Solution of the Joints Problem

This solution was first produced by Guth-Katz for the three-dimensional case in [GK10] and later extended to the general case by Quilodrán in [Qui10] and independently at the same time by Kaplan-Sharir-Shustin in [KSS10].

Theorem 4.2.1. In \mathbb{R}^n we have

$$J(L) \lesssim_n L^{\frac{n}{n-1}}.$$

We begin with the fundamental lemma to this proof. The key extension of the polynomial method here over that of Chapter 2 is the examination of directional derivatives of polynomials.

Lemma 4.2.2. If \mathcal{L} is a set of lines in \mathbb{R}^n that determines J joints, then one of the lines contains at most $nJ^{\frac{1}{n}}$ joints.

Proof. Let P be a non-zero polynomial that vanishes at every joint of \mathcal{L} and assume that the degree of P is minimal. By parameter counting (Lemma 2.2.2) the degree of P is at most $nJ^{\frac{1}{n}}$. (To see this, set $D = \lfloor nJ^{\frac{1}{n}} \rfloor$ and then notice that $J < \binom{D+n}{n}$.)

We proceed by contradiction. Assume every line contains more than $nJ^{\frac{1}{n}}$ joints. Now P must vanish on every line in \mathcal{L} as the degree of P is less than the number of joints contained in the line, which are points of intersection between the line and $Z(P)$.

We now examine the gradient of P at each joint in \mathcal{L} . We will need a fact about gradients for this, which we encapsulate in the following lemma for clarity.

Lemma 4.2.3. *If x is a joint of \mathcal{L} , and if a smooth function $F : \mathbb{R}^n \rightarrow \mathbb{R}$ vanishes on the lines of \mathcal{L} , then ∇F vanishes at x .*

Proof. The joint x is contained in n non-coplanar lines l_1, \dots, l_n with directions v_1, \dots, v_n respectively. Now consider the directional derivative for a particular v_i

$$\frac{\partial F}{\partial v_i}(x) = \lim_{t \rightarrow 0} \frac{\overbrace{F(x + tv_i)}^{F \equiv 0 \text{ on a line in } \mathcal{L}} - \overbrace{F(x)}^{F \equiv 0 \text{ on joints}}}{t} = \frac{0}{t} = 0.$$

Notice that $\frac{\partial F}{\partial v_i} = \langle \nabla F, v_i \rangle$ so since we have this for each v_i , and the set of v_i 's form a basis of \mathbb{R}^n we have that $\nabla F(x) = 0$. \square

So we see that the partial derivatives of P vanish at each joint. The derivatives are polynomials of smaller degree than P and since P was assumed to be a non-zero polynomial of minimal degree that vanishes at each joint, each derivative of P must be identically zero. This implies P must be constant, which is a contradiction. \square

Finally we can prove the main result.

Proof. Lemma 4.2.2 tells us that if we remove a line from our collection, we are removing at most $nJ(L)^{\frac{1}{n}}$ joints. By repeating this process, we get the chain of inequalities

$$\begin{aligned} J(L) &\leq J(L-1) + n(J(L))^{\frac{1}{n}} \\ &\leq J(L-2) + 2 \left[n(J(L))^{\frac{1}{n}} \right] \\ &\leq J(L-3) + 3 \left[n(J(L))^{\frac{1}{n}} \right] \\ &\vdots \\ &\leq L \left[n(J(L))^{\frac{1}{n}} \right]. \end{aligned}$$

Rearranging we have

$$J(L)^{\frac{n-1}{n}} \lesssim_n L \implies J(L) \lesssim_n L^{\frac{n}{n-1}}.$$

\square

Chapter 5

Szemerédi-Trotter Theorem

Incidence Geometry is the quantitative study of incidence relations between simple geometric objects, such as lines or low degree curves. In this chapter we will study the application of the polynomial method to incidence geometry by proving a fundamental theorem in the field. We have already seen an incidence problem in the previous chapter on the Joints problem. By developing the powerful tool of polynomial partitioning we shall see the key role that the topology of \mathbb{R}^n can play in such problems.

The Szemerédi-Trotter theorem is a fundamental result in the field of Incidence Geometry, originally proved by an involved cell decomposition argument of Szemerédi-Trotter in [ST83] and later given a shorter proof using the crossing-number inequality by Székely in [Szé97].

Theorem 5.0.1 (Szemerédi-Trotter). *Let $\mathcal{S} \subset \mathbb{R}^2$ be a finite set of points and let \mathcal{L} be a finite set of lines in \mathbb{R}^2 . We define*

$$I(\mathcal{S}, \mathcal{L}) = \{(p, \ell) \in \mathcal{S} \times \mathcal{L} \mid p \in \ell\}$$

to be the set of incidences between \mathcal{S} and \mathcal{L} .

Then

$$|I(\mathcal{S}, \mathcal{L})| \lesssim (|\mathcal{S}||\mathcal{L}|)^{2/3} + |\mathcal{S}| + |\mathcal{L}|.$$

5.1 The Trivial Bounds

In planar geometry, we have the following dual statements: two points determine at least one line and every pair of lines intersect in at most one point. Using this we can prove the following bounds on $I(\mathcal{S}, \mathcal{L})$:

Lemma 5.1.1 (Trivial Bounds). *For a set of points \mathcal{S} and lines \mathcal{L} we have*

$$|I(\mathcal{S}, \mathcal{L})| \leq |\mathcal{S}|^2 + |\mathcal{L}|,$$

and

$$|I(\mathcal{S}, \mathcal{L})| \leq |\mathcal{L}|^2 + |\mathcal{S}|.$$

Proof. To see this, count the lines that have at most one point in P on them. These contribute at most $|\mathcal{L}|$ incidences. Every other line has at least two points in \mathcal{S} . The total number of incidences on these lines is at most $|\mathcal{S}|^2$ as otherwise by the pigeonhole principle there would exist a $p \in \mathcal{S}$ that lies in over $|\mathcal{S}|$ lines, and each of these lines would have an additional point on it. This would imply there are more than $|\mathcal{S}|$ points, a contradiction.

Interchanging the roles of \mathcal{L} and \mathcal{S} achieves the other bound as two lines intersect in at most one point. \square

By double counting we can achieve a slightly less trivial bound, which as we will see in the next section turns out to be tight in finite fields.

Lemma 5.1.2 (Double Counting Bounds). *For a set of points \mathcal{S} and lines \mathcal{L} we have*

$$|I(\mathcal{S}, \mathcal{L})| \lesssim |\mathcal{S}||\mathcal{L}|^{\frac{1}{2}} + |\mathcal{L}|,$$

and

$$|I(\mathcal{S}, \mathcal{L})| \lesssim |\mathcal{L}||\mathcal{S}|^{\frac{1}{2}} + |\mathcal{S}|.$$

Proof. We bound the number of incidences using the Cauchy-Schwarz inequality followed by double counting.

$$|I(\mathcal{S}, \mathcal{L})|^2 = \left(\sum_{\ell \in \mathcal{L}} \sum_{p \in \mathcal{S}} \mathbb{1}[p \in \ell] \right)^2$$

Applying the Cauchy-Schwarz inequality on the collection \mathcal{L} we attain

$$\begin{aligned} &\leq |\mathcal{L}| \cdot \sum_{\ell \in \mathcal{L}} \left(\sum_{p \in \mathcal{S}} \mathbb{1}[p \in \ell] \right)^2 \\ &= |\mathcal{L}| \cdot \sum_{p_1, p_2 \in \mathcal{S}} \sum_{\ell \in \mathcal{L}} \mathbb{1}[p_1 \in \ell] \mathbb{1}[p_2 \in \ell] \\ &= |\mathcal{L}| \left(\sum_{p_1 = p_2 \in \mathcal{S}} \sum_{\ell \in \mathcal{L}} \mathbb{1}[p_1 \in \ell] \right. \\ &\quad \left. + \sum_{p_1 \neq p_2 \in \mathcal{S}} \sum_{\ell \in \mathcal{L}} \mathbb{1}[p_1 \in \ell] \mathbb{1}[p_2 \in \ell] \right). \end{aligned}$$

The diagonal terms above give a contribution of $|I(\mathcal{S}, \mathcal{L})|$ by definition. The off diagonal terms can be trivially bounded by $|\mathcal{S}|^2$. Hence

$$|I(\mathcal{S}, \mathcal{L})|^2 \leq |\mathcal{L}|(|I(\mathcal{S}, \mathcal{L})| + |\mathcal{S}|^2).$$

Using Lemma 5.1.1 to bound the $|I(\mathcal{S}, \mathcal{L})|$ term we achieve

$$|I(\mathcal{S}, \mathcal{L})|^2 \leq |\mathcal{L}|^2 + 2|\mathcal{L}||\mathcal{S}|^2.$$

This implies

$$I(\mathcal{S}, \mathcal{L}) \lesssim |\mathcal{S}||\mathcal{L}|^{\frac{1}{2}} + |\mathcal{L}|.$$

Repeating the above proof interchanging the roles \mathcal{S} and \mathcal{L} achieves the other bound. \square

5.2 Examples

The following example shows we can not improve beyond the double counting bounds shown in Lemma 5.1.2 in a finite field \mathbb{F}^2 .

Example 5.1 (Finite Fields). *Consider the set of points $\mathcal{S} = \mathbb{F}^2$ and let \mathcal{L} be the set of all lines in \mathbb{F}^2 . Every line contains exactly $|\mathbb{F}|$ many points of \mathcal{S} , so we have $|\mathbb{F}|^3$ incidences. So both sides of the double counting bounds (Lemma 5.1.2) are comparable:*

$$I(\mathcal{S}, \mathcal{L}) = |\mathbb{F}|^3 \sim (|\mathbb{F}|^2)(|\mathbb{F}|^2)^{1/2} + |\mathbb{F}|^2.$$

In contrast, the following examples turn out to be the best possible over \mathbb{R} . We will later prove that these are optimal for the Szemerédi–Trotter Theorem. We denote a line in \mathbb{R}^2 as follows

$$\ell_{m,c} = \{(x, y) \in \mathbb{R}^2 \mid y = mx + c\}.$$

First we examine two trivial examples that explain the necessity of the $|\mathcal{S}| + |\mathcal{L}|$ terms in Theorem 5.0.1.

Example 5.2. *Let $|\mathcal{S}| = 1$ and suppose all $|\mathcal{L}|$ lines in \mathcal{L} pass through this point. Then both sides of the Szemerédi–Trotter (Theorem 5.0.1) inequality are comparable as*

$$|I(\mathcal{S}, \mathcal{L})| = |\mathcal{L}| \sim 1 + |\mathcal{L}| + |\mathcal{L}|^{\frac{2}{3}}.$$

Consider the dual case where $|\mathcal{L}| = 1$ and suppose all $|\mathcal{S}|$ points in \mathcal{S} lie on \mathcal{L} . Again we have

$$|I(\mathcal{S}, \mathcal{L})| = |\mathcal{S}| \sim 1 + |\mathcal{S}| + |\mathcal{S}|^{\frac{2}{3}}.$$

The next two examples deal with the more general case, where the $(|\mathcal{S}||\mathcal{L}|)^{2/3}$ term dominates Theorem 5.0.1.

Example 5.3. Let N be a large integer and consider the following collections in \mathbb{R}^2

$$\mathcal{S} = \{(a, b) \in \mathbb{Z}^2 \mid a \in [1, N], b \in [1, 2N^2]\},$$

$$\mathcal{L} = \{\ell_{m,c} \in \mathbb{R}^2 \mid m, c \in \mathbb{Z}, m \in [1, N], c \in [1, N^2]\}.$$

The collection \mathcal{S} contains $2N^3$ points and \mathcal{L} contains N^3 lines. Every line in \mathcal{L} contains N points in \mathcal{S} as for each $x \in [1, N]$ the y -coordinate of $\ell_{m,c}$, $mx + c$, gives a different integer in $[1, 2N^2]$. Hence there are N^4 incidences. Both sides of the Szemerédi-Trotter (Theorem 5.0.1) inequality are comparable as

$$|I(\mathcal{S}, \mathcal{L})| = N^4 \sim (N^3)^{\frac{2}{3}}(N^3)^{\frac{2}{3}} \sim |\mathcal{S}|^{2/3}|\mathcal{L}|^{2/3} + |\mathcal{S}| + |\mathcal{L}|.$$




Figure 5.1: Example 5.3 in the case $N = 3$. (with the x, y axis reversed)

Example 5.4. Let $N > 1$ be a large even integer and let $1 < R \lesssim N$ be another integer. Consider the collections in \mathbb{R}^2 given by

$$\mathcal{S} = \{(a, b) \in \mathbb{Z}^2 \mid (a, b) \in [-N/2, N/2] \times [-N/2, N/2]\},$$

$$\mathcal{L} = \{\ell \mid \ell \text{ contains between } R \text{ and } 2R \text{ points of } \mathcal{S}\}.$$

We begin by estimating how many lines of \mathcal{L} pass through a given point of the regular grid \mathcal{S} . Let $\ell \in \mathcal{L}$ and $p \in \mathcal{S}$. The closest point $p' \in \mathcal{S}$ such that $p \neq p'$ and $p' \in \ell$ must lie in a square centred at p of side length $\sim N/R$. This follows from the fact that there are at least $\sim R$ points of \mathcal{S} in ℓ and hence the projections of these points to the axes can be separated by at most $\sim N/R$. Taking each possible combination of these we can conclude that there are $\lesssim N^2/R^2$ lines in \mathcal{L} through a given point p .

We now claim that there are $\gtrsim N^2/R^2$ distinct such lines. We need only consider the points in the upper right quadrant of \mathcal{S} as the problem is symmetrical. Further, we restrict ourselves to considering lines with slopes m satisfying $\frac{1}{2} < m < 2$. For such a line to contain R points of \mathcal{S} it is enough that $m = \frac{l}{k} \in \mathbb{Q}$ with $\gcd(l, k) = 1$ and $l, k \in [\frac{N}{2R}, \frac{N}{R}]$. There are $\gtrsim N^2/R^2$ such pairs (l, k) , as the proportion of pairs that share a factor of 2 is $\frac{1}{2^2}$ and the proportion of pairs that share a factor of 3 is $\frac{1}{3^2}$, and in general the proportion that shares a factor of k is $\frac{1}{k^2}$. We have that $\sum_{k>1} \frac{1}{k^2} < \frac{2}{3} < 1$ and hence there are $\gtrsim N^2/R^2$ distinct lines in \mathcal{L} through each point. Taking account of what we have shown

$$\begin{aligned} |\mathcal{S}| &\sim N^2, \\ |\mathcal{L}| &\sim |\mathcal{S}| \frac{N^2}{R^2} \frac{1}{R} \sim \frac{N^4}{R^3}, \\ |I(\mathcal{S}, \mathcal{L})| &\sim |\mathcal{S}| \frac{N^2}{R^2} \sim \frac{N^4}{R^2}. \end{aligned}$$

We can see that both sides of the Szemerédi-Trotter inequality (Theorem 5.0.1) are comparable as

$$|I(\mathcal{S}, \mathcal{L})| \sim \frac{N^4}{R^2} \sim (N^2)^{\frac{2}{3}} \left(\frac{N^4}{R^3} \right)^{\frac{2}{3}} \sim |\mathcal{S}|^{\frac{2}{3}} |\mathcal{L}|^{\frac{2}{3}} + |\mathcal{S}| + |\mathcal{L}|.$$

5.3 Ham Sandwich Theorems

The above examples suggest that the topology of \mathbb{R}^2 plays a key role in this incidence problem. We shall now introduce the method of polynomial partitioning, which can be seen as the topological analogue to the vanishing lemma (Lemma 2.2.3) we used in the previous chapters.

The topological input in the proof will be provided by the following well-known result of Borsuk-Ulam. Let \mathbb{S}^n denote the unit n -sphere in \mathbb{R}^{n+1} .

Theorem 5.3.1 (Borsuk-Ulam). *A map ϕ is said to be antipodal if it obeys $\phi(-x) = -\phi(x)$ for all x in its domain. Suppose $\phi : \mathbb{S}^N \rightarrow \mathbb{R}^N$ is a continuous antipodal mapping. Then the image of ϕ contains 0.*

In the appendix to this manuscript (see Appendix A.0.1), we present a beautiful combinatorial proof from Matousek's book *Using the Borsuk-Ulam Theorem* (see [Mat03]). The proof proceeds by constructively proving the Tucker Lemma, and then using this to imply the Borsuk-Ulam theorem.

Let us now define some useful definitions going forward.

Definition 5.3.1 (Bisection of a Set). A function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is said to bisect an open

set U with volume¹ $\text{Vol}(U) < \infty$ if

$$\text{Vol}\{x \in U \mid f(x) > 0\} = \text{Vol}\{x \in U \mid f(x) < 0\} = \frac{1}{2}\text{Vol}(U).$$

Analogously, a function f is said to bisect a finite set S if both

$$|\{x \in S \mid f(x) > 0\}| \leq \frac{|S|}{2},$$

and

$$|\{x \in S \mid f(x) < 0\}| \leq \frac{|S|}{2}.$$

We now show that as a consequence of Theorem 5.3.1, it is always possible to find a non-zero continuous function that bisects every open set in a collection.

Theorem 5.3.2 (General Ham Sandwich Theorem). *Let V be a finite dimensional vector space of continuous functions $f : \mathbb{R}^n \rightarrow \mathbb{R}$ such that for any non-zero function f , $Z(f)$ has zero Lebesgue measure. Let $U_1, U_2, \dots, U_N \subset \mathbb{R}^n$ be finite volume open sets with $N < \dim V$.*

Then there exists a non-zero function $f \in V$ that bisects each U_i .

Proof. Define the functions $\{\phi_i\}_{i=1}^N$, $\phi_i : V \setminus \{0\} \rightarrow \mathbb{R}$ by:

$$\phi_i(f) = \text{Vol}(\{x \in U_i \mid f(x) > 0\}) - \text{Vol}(\{x \in U_i \mid f(x) < 0\}).$$

Since $Z(f)$ has measure zero, it is easy to see that $\phi_i(f) = 0$ if and only if f bisects U_i . Notice also that $\phi_i(-f) = -\phi_i(f)$, hence ϕ_i is antipodal.

We now show each ϕ_i is continuous. It is enough to show that if U is a finite volume open set, then the measure of $\{x \in U \mid f(x) > 0\}$ depends continuously on $f \in V \setminus \{0\}$.

Suppose $f_n \rightarrow f$ in V for some $f, f_n \in V \setminus \{0\}$. f_n converges to f in the topology of V , so it follows it must converge pointwise. Pick any $\varepsilon > 0$. By Egorov's theorem, we can find a subset $E \subset U$ so that $f_n \rightarrow f$ uniformly pointwise on $U \setminus E$ with $m(E) < \varepsilon$. By hypothesis, $m(Z(f)) = 0$ and $m(U) < \infty$. Since the Lebesgue measure is continuous we can choose δ such that $m(\{x \in U \mid |f(x)| < \delta\}) < \varepsilon$.

Now we choose n sufficiently large that $|f_n(x) - f(x)| < \delta$ on $U \setminus E$. Then we have

$$\begin{aligned} & |m(\{x \in U \mid f_n(x) > 0\}) - m(\{x \in U \mid f(x) > 0\})| \\ &= |m(\{x \in U \mid f_n(x) > 0, f(x) > 0\}) + m(\{x \in U \mid f_n(x) > 0, f(x) < 0\}) \\ &\quad - m(\{x \in U \mid f(x) > 0, f_n(x) > 0\}) - m(\{x \in U \mid f(x) > 0, f_n(x) < 0\})| \\ &= |m(\{x \in U \mid f_n(x) > 0, f(x) < 0\}) - m(\{x \in U \mid f(x) > 0, f_n(x) < 0\})|. \end{aligned}$$

¹Volume is given by the Lebesgue measure.

By the triangle inequality we have

$$\leq m(\{x \in U \mid f_n(x) > 0, f(x) < 0\}) + m(\{x \in U \mid f(x) > 0, f_n(x) < 0\}).$$

The remaining two terms are symmetrical, so we proceed by bounding $m(\{x \in U \mid f(x) > 0, f_n(x) < 0\})$. We can write

$$\begin{aligned} m(\{x \in U \mid f(x) > 0, f_n(x) < 0\}) &= m(\{x \in U \mid f(x) > 0, f_n(x) < 0\} \cap E) \\ &\quad + m(\{x \in U \mid f(x) > 0, f_n(x) < 0\} \setminus E) \\ &< \varepsilon + m(\{x \in U \mid f(x) > 0, f_n(x) < 0\} \setminus E). \end{aligned}$$

We claim that $\{x \in U \mid f(x) > 0, f_n(x) < 0\} \setminus E \subset \{x \in U \mid |f(x)| < \delta\} \setminus E$. Notice that if $f(x) > 0$ and $f_n(x) < 0$ we have

$$|f(x)| = f(x) = f_n(x) + f(x) - f_n(x) < \delta$$

and similarly that

$$|f_n(x)| = -f_n(x) = +f(x) - f_n(x) - f(x) < \delta.$$

This establishes our claim and we can conclude that

$$m(\{x \in U \mid f(x) > 0, f_n(x) < 0\}) < 2\varepsilon.$$

Since ε was arbitrary each ϕ_i is continuous.

We now combine each ϕ_i into the map $\phi : V \setminus \{0\} \rightarrow \mathbb{R}^N$ by $\phi = (\phi_1, \dots, \phi_N)$. Since $\dim V > N$, select a subspace $W < V$ such that $\dim W = N + 1$. Now choose an isomorphism of W with \mathbb{R}^{N+1} and notice that since for all $\lambda > 0$ we have $\phi(\lambda x) = \phi(x)$ and hence that ϕ is simply a function on the unit sphere. Now the map $\phi : \mathbb{S}^N \rightarrow \mathbb{R}^N$ is antipodal and continuous. By the Borsuk-Ulam Theorem, there exists an $f \in \mathbb{S}^N \subset V \setminus \{0\}$ such that $\phi(f) = 0$. \square

We are now going to use the Ham Sandwich Theorem to extend the conclusion to the bisection of finite sets also. The main idea will be to replace each of our points in the finite set with δ -balls and apply the previous theorem.

Corollary 5.3.1 (Finite Ham Sandwich Theorem). *Let S_1, \dots, S_N be finite sets in \mathbb{R}^n and let D be such that $N < \binom{D+n}{n}$. Then there exists a non-zero $P \in \text{Poly}_D(\mathbb{R}^n)$ that bisects each S_i .*

Proof. Let us equip the space $\text{Poly}_D(\mathbb{R}^n)$ with the L^1 norm. For each $\delta > 0$, define $U_{i,\delta}$ to be the union of δ -balls centred at the points of S_i . By Theorem 5.3.2, we can find a non-zero P_δ with degree at most D that bisects each $U_{i,\delta}$. By rescaling we can assume

$P_\delta \in \mathbb{S}^M \subset \text{Poly}_D(\mathbb{R}^n) \setminus \{0\}$. Since \mathbb{S}^M compact, we can find a sequence $\delta_m \rightarrow 0$ so that P_{δ_m} converges to P in \mathbb{S}^M . Since the coefficients of P_{δ_m} converge to those of P , P_{δ_m} converges to P uniformly on compact sets.

We claim P bisects each S_i . By contradiction, suppose $P > 0$ on more than half the points of S_i , say on the points of S_i^+ . Choosing ε sufficiently small, we can assume $P > 0$ on the ε -ball around each point of S_i^+ . Further, we can choose ε such that each ε -ball is disjoint. Since P_{δ_m} converges uniformly, we can find m sufficiently large such that $P_{\delta_m} > 0$ on the ε -ball around each point of S_i^+ . By making m large, we can also arrange that $\delta_m < \varepsilon$. Thus $P_{\delta_m} > 0$ on more than half the volume of the set U_{i,δ_m} , a contradiction. \square

We now utilise the above results to prove a powerful cell decomposition technique known as polynomial partitioning.

Theorem 5.3.3 (Polynomial Partitioning). *If S is a finite subset of \mathbb{R}^n and D any degree, there exists a non-zero polynomial P of degree at most D such that $\mathbb{R}^n \setminus Z(P)$ is a disjoint union of $\sim D^n$ open sets O_i referred to as cells each containing $\lesssim_n |S|D^{-n}$ points.*

Proof. The main idea of the proof is the repeated application of the Finite Ham Sandwich Theorem. We begin by finding a polynomial P_1 of degree 1 that bisects S . This partitions $\mathbb{R} \setminus Z(P_1)$ into two disjoint open sets according to the sign of P_1 ,

$$P_1^+ = \{x \in S \mid P_1(x) > 0\}$$

and

$$P_1^- = \{x \in S \mid P_1(x) < 0\}$$

each containing at most $|S|/2$ points. We then bisect both of these sets using another polynomial P_2 . There are four sign conditions on P_1 and P_2 , these being the four possible intersections of the sets P_1^\pm and P_2^\pm , and the subset for each sign condition contains at most $|S|/4$ points of S . Continuing this process to define polynomials P_3, P_4, \dots , where the polynomial P_j simultaneously bisects 2^{j-1} finite sets. By the Finite Ham Sandwich Theorem, we can find a polynomial of degree $\lesssim 2^{j/n}$ that simultaneously bisects 2^{j-1} finite sets.

If we repeat this procedure J times and defining $P = \prod_{i=1}^J P_i$, then $\mathbb{R}^n \setminus Z(P)$ is the disjoint union of 2^J open sets each containing $\leq |S|2^{-J}$ points of S . Now we choose the largest J such that $\deg(P) < D$ which is equivalent to $\sum_{j=0}^J c(n)2^{j/n} \leq D$. But $\sum_{j=0}^J 2^{j/n}$ is a geometric series so we can arrange that $\deg(P) < D$ for $D \leq c(n)2^{J/n}$. The number of points in each O_i is $\leq |S|2^{-J} \leq c(n)|S|D^{-n}$. \square

There is a crucial point to note about polynomial partitioning. The above theorem does not guarantee anything about the distribution of points between $Z(P)$ and its complement.

This is made most clear looking at the following two extremes. If all points lie in the complement of $Z(P)$ then we have an optimal equidistribution of points, and can often use trivial bounds in a divide-and-conquer style argument. On the other hand, in the case all points are contained in $Z(P)$ we have many points in an algebraic surface of controlled degree, so we can try and use tools from algebraic geometry. Generally there will be some points in both $Z(P)$ and its complement which we need to deal with separately.

5.4 Proof of the Szemerédi-Trotter Theorem

We now can prove the Szemerédi-Trotter Theorem (Theorem 5.0.1) using polynomial partitioning.

Proof of the Szemerédi-Trotter Theorem. Let $|\mathcal{S}| = S$ and $|\mathcal{L}| = L$.

Considering the regime where $L > S^2$ and applying the trivial bounds (Lemma 5.1.1) yields

$$|I(\mathcal{S}, \mathcal{L})| \leq S^2 + L \sim L \sim (SL)^{2/3} + S + L.$$

Similarly in the case $L < S^{\frac{1}{2}}$ we have

$$|I(\mathcal{S}, \mathcal{L})| \leq L^2 + S \sim S \sim (SL)^{2/3} + S + L.$$

Hence we need only consider the case $S^{\frac{1}{2}} \leq L \leq S^2$.

Let D be a fixed degree to be chosen later. By Theorem 5.3.3, there exists a polynomial P of degree D such that $\mathbb{R}^2 \setminus Z(P)$ splits into $\sim D^2$ disjoint open sets each containing $\lesssim SD^{-2}$ points of \mathcal{S} . Let $\{O_i \mid i \in \Pi\}$ denote the family of the cells in the decomposition. Let $\mathcal{S}_i = \mathcal{S} \cap O_i$ and \mathcal{L}_i denote the lines that intersect the interior of each O_i respectively. We define the following pairs of complementary sets

$$\mathcal{S}_c = \{x \in \mathcal{S} \mid x \notin Z(P)\}$$

$$\mathcal{S}_z = \{x \in \mathcal{S} \mid x \in Z(P)\}$$

$$\mathcal{L}_c = \{\ell \in \mathcal{L} \mid \ell \not\subset Z(P)\}$$

$$\mathcal{L}_z = \{\ell \in \mathcal{L} \mid \ell \subset Z(P)\}$$

Note that $\mathcal{S} = \mathcal{S}_c \cup \mathcal{S}_z$, $\mathcal{L} = \mathcal{L}_c \cup \mathcal{L}_z$. We can now write our total line-point incidences as the following sum

$$|I(\mathcal{S}, \mathcal{L})| = |I(\mathcal{S}_c, \mathcal{L})| + |I(\mathcal{S}_z, \mathcal{L}_z)| + |I(\mathcal{S}_z, \mathcal{L}_c)|.$$

We begin by examining the $I(\mathcal{S}_c, \mathcal{L})$ term

$$|I(\mathcal{S}_c, \mathcal{L})| = \sum_{i \in \Pi} |I(\mathcal{S}_i, \mathcal{L}_i)|.$$

Using our trivial bound (Lemma 5.1.1) in each cell we attain

$$\leq \sum_{i \in \Pi} \mathcal{S}_i^2 + \sum_{i \in \Pi} \mathcal{L}_i.$$

If a line ℓ is not contained entirely in $Z(P)$ then it can intersect $Z(P)$ at most D times, so each line intersects at most $D + 1$ cells. Hence $\sum_{i \in \Pi} \mathcal{L}_i \leq (D + 1)L$. So we have

$$\begin{aligned} |I(\mathcal{S}_c, \mathcal{L})| &\lesssim LD + SD^{-2} \sum_{i \in \Pi} \mathcal{S}_i \\ &\leq LD + S^2 D^{-2}. \end{aligned}$$

The number of lines in \mathcal{L}_z is at most D as $Z(P)$ can contain at most D degree 1 factors. So we have by our trivial bounds (Lemma 5.1.1)

$$|I(\mathcal{S}_z, \mathcal{L}_z)| \leq S + D^2.$$

Each line in \mathcal{L}_c has at most D intersection points with $Z(P)$ so it has at most D incidences with \mathcal{S}_z . Hence

$$|I(\mathcal{S}_z, \mathcal{L}_c)| \leq LD.$$

Together we have now

$$|I(\mathcal{S}, \mathcal{L})| \lesssim LD + S^2 D^{-2} + S + D^2.$$

We optimise $LD + S^2 D^{-2}$ by choosing D such that both terms comparable to each other and hence $D \sim S^{\frac{2}{3}} L^{-\frac{1}{3}}$. From our restriction $S^{\frac{1}{2}} \leq L \leq S^2$ we have $S^{\frac{2}{3}} L^{-\frac{1}{3}} \geq 1$ and $D^2 \sim S^{\frac{4}{3}} L^{-\frac{2}{3}} \leq S$, so we achieve

$$|I(\mathcal{S}, \mathcal{L})| \lesssim (SL)^{2/3} + S.$$

Finally combining this with our result for the regime where $L > S^2$ we achieve the full Szemerédi-Trotter inequality

$$|I(\mathcal{S}, \mathcal{L})| \lesssim (SL)^{2/3} + S + L.$$

□

There are two key things to note about the above proof. First, the key role that the topology of \mathbb{R}^2 plays. We use topology twice in the above proof. Firstly by using polynomial partitioning, which itself is a consequence of the Borsuk-Ulam theorem from topology. We also appeal to topology in our claim that each line intersects at most $\deg P + 1$ cells of $\mathbb{R}^2 \setminus Z(P)$. It is a worthwhile heuristic to develop that polynomial partitioning may be useful for incidence problems where the best examples in a finite field (which is only equipped with the trivial topology) do not coincide with the best known examples over the

reals. Secondly, the above proof illustrates the surprising power of polynomial partitioning. We were able to use very trivial bounds in each cell to achieve an asymptotically tight overall bound.

Chapter 6

Counting Circle Tangencies

Here we shall discuss the problem of counting the number of tangencies in a suitably non-degenerate collection of circles. We say two circles are tangent if their intersection contains a single point. The set of pairs of circles in a collection \mathcal{C} which are mutually tangent are called the tangencies and the collection of all tangencies is denoted $\tau(\mathcal{C})$.

Lemma 6.0.1 (Trivial Bound). *Let \mathcal{C} be an arbitrary finite collection of circles. Then the number of tangencies $|\tau(\mathcal{C})|$ is bounded by*

$$|\tau(\mathcal{C})| \leq |\mathcal{C}|^2.$$

Proof. For each pair of circles in \mathcal{C} there can be at most one tangency between them. Since there are $\leq |\mathcal{C}|^2$ pairs, the bound follows. \square

Stated for arbitrary collections of circles the problem is not particularly interesting as the above bound turns out to be asymptotically tight, as the following example shows.

Example 6.1. *Denote a circle centred at (x, y) with radius r as $\gamma(x, y, r)$. Consider the following collection of $N + 1$ circles:*

$$\begin{aligned} C_0 &= \gamma(0, 0, 1) \\ C_1 &= \gamma\left(\frac{1}{2}, 0, \frac{1}{2}\right) \\ C_2 &= \gamma\left(\frac{3}{4}, 0, \frac{1}{4}\right) \\ &\vdots \\ C_N &= \gamma\left(1 - \frac{1}{2^N}, 0, \frac{1}{2^N}\right) \end{aligned}$$

Each circle in our collection $\mathcal{C} = \{C_i \mid 0 \leq i \leq N\}$ is tangent to N other circles at the point $(1, 0)$. Hence $|\tau(\mathcal{C})| \sim N^2 \sim |\mathcal{C}|^2$.

In light of this example, in order to obtain a non-trivial example we shall look at collections of circles that satisfy a non-degeneracy condition. We consider collections of




Figure 6.1: A collection of circles in \mathbb{R}^2 with N^2 tangencies.

circles such that no three circles are mutually tangent at a common point. The best known examples for collections that satisfy this non-degeneracy condition leverage our asymptotically tight Szemerédi-Trotter Theorem (Theorem 5.0.1) examples:

Example 6.2. *Let \mathcal{S} , \mathcal{L} be collections of N points and N lines respectively such that the Szemerédi-Trotter Theorem’s bound is asymptotically tight. In particular, they determine $\sim N^{4/3}$ point-line incidences. Let C_1 denote the collection of unit circles centred at the points of \mathcal{S} . Let C_2 denote the collection of lines obtained by translating each line $\ell \in \mathcal{L}$ one unit in the ℓ^\perp direction. If $(p, \ell) \in \mathcal{S} \times \mathcal{L}$ is a point-line incidence from our original collection, then the corresponding circle-line pair will be tangent. Performing an inversion transform about a point that does not lie in any of the circles or lines, our collection $\mathcal{C} = C_1 \cup C_2$ becomes a collection of $2N$ circles that determine at least $N^{4/3}$ tangencies.*

We shall now present a bound from a recent paper of Ellenberg-Solymosi-Zahl which uses the polynomial method (See [ESZ16]).

Theorem 6.0.2 ([ESZ16]). *Given a finite collection \mathcal{C} of N circles in the plane such that no three are mutually tangent at a common point, then the number of tangencies $|\tau(\mathcal{C})|$ obeys*

$$|\tau(\mathcal{C})| \lesssim N^{3/2}.$$

It is not obvious how to apply the polynomial method in the current formulation. We will perform a lifting of our circles into algebraic curves in \mathbb{R}^3 , transforming tangencies between circles in the plane to incidences between their lifted curves in \mathbb{R}^3 . This transforms the problem from a tangency problem to an incidence problem which can therefore be attacked with the tools available to incidence geometry. As an additional benefit to this lifting, it reduces the required degree of polynomial needed to interpolate the incidences. In general Lemma 2.2.2 tells us that if we are to interpolate a set of M points in \mathbb{R}^d

by a polynomial P , then the degree of P is $O(M^{1/d})$, so increasing the dimension of the problem yields a lower degree of the polynomial required.

6.1 Lifting of circles to \mathbb{R}^3

We shall now discuss this lifting in detail. Let $\gamma \subset \mathbb{R}^2$ be a circle in the plane of radius r_γ centred at (x_γ, y_γ) . We define the lifting transform of the circle as

$$\beta(\gamma) = \left\{ (x, y, z) \in \mathbb{R}^3 \mid (x - x_\gamma)^2 + (y - y_\gamma)^2 = r_\gamma^2, \ z = -\frac{x - x_\gamma}{y - y_\gamma} \right\}.$$

Clearly $\beta(\gamma)$ is an algebraic curve, so now we shall examine what happens to pairs of tangent circles under this transformation. Notice that z is defined as the slope of the tangent line at the point $(x, y) \in \gamma$.

Lemma 6.1.1. *Let β be the transform defined as above. Then two circles $\gamma, \gamma' \subset \mathbb{R}^2$ are tangent if and only if $\beta(\gamma) \cap \beta(\gamma') \neq \emptyset$.*

Proof. If γ and γ' are tangent then there exists a point $(x, y) \in \gamma \cap \gamma'$ and at a point of tangency we have that the slopes of the tangent lines coincide for γ and γ' . Explicitly $-\frac{x - x_\gamma}{y - y_\gamma} = -\frac{x - x_{\gamma'}}{y - y_{\gamma'}} = z$. Hence $(x, y, z) \in \beta(\gamma) \cap \beta(\gamma')$ and thus $\beta(\gamma) \cap \beta(\gamma') \neq \emptyset$.

In the other direction, assume there exists some $(x, y, z) \in \beta(\gamma) \cap \beta(\gamma')$. Clearly $(x, y) \in \gamma \cap \gamma'$ and the slopes of the tangent lines at this point are equal, hence we conclude γ is tangent to γ' . \square

This one-to-one correspondence between tangencies in \mathbb{R}^2 and incidences in \mathbb{R}^3 is the key idea behind the proof. It may be viewed as the graph of a Gaussian map in a certain coordinate system.

In the course of the proof we will need a Bézout-type result which bounds the number of intersections between a curve and a surface with no common components.

Lemma 6.1.2 (Bézout's Lemma in \mathbb{R}^3). *Let $Z(P)$ be a surface over $\mathbb{R}[X, Y, Z]$ and $Z(Q_1, Q_2)$ be a curve over $\mathbb{R}[X, Y, Z]$. Suppose that P, Q_1, Q_2 have no pairwise common factors. Then we can bound the cardinality of the intersection of the zero sets by*

$$|Z(P) \cap Z(Q_1, Q_2)| \leq \deg P \deg Q_1 \deg Q_2.$$

This result follows from Theorem 12.3 in Fulton's book "*Intersection Theory*" (See [Ful98], page 223).

6.2 Ellenberg-Solymosi-Zahl's Proof of Theorem 6.0.2

In this proof, we will need to ensure our tangencies are sufficiently uniformly distributed among our circles. We can refine our collection such that this is the case by the following lemma.

Lemma 6.2.1 (Uniform Refinement). *Let \mathcal{C} be a collection of N circles and suppose that $|\tau(\mathcal{C})| \gtrsim N^\alpha$. Then we can refine our collection to a collection $\mathcal{C}' \subset \mathcal{C}$ such that every circle in \mathcal{C}' is tangent to $\gtrsim N^{\alpha-1}$ other circles in \mathcal{C}' , $|\tau(\mathcal{C}')| \gtrsim N^\alpha$, and $|\mathcal{C}'| \gtrsim N^{\alpha/2}$.*

Proof. We proceed by a stopping-time argument. Let $\tau(\mathcal{C})$ be the set of tangencies of the circles in \mathcal{C} and let c_0 be a constant such that $|\tau(\mathcal{C})| \geq c_0 N^\alpha$. Let $\mathcal{C}_0 = \mathcal{C}$ and let c_1 be a fixed constant to be chosen later. If there exists a circle $\gamma \in \mathcal{C}_0$ such that $|\{\gamma' \mid (\gamma, \gamma') \in \tau(\mathcal{C}_0)\}| < c_1 N^{1/2}$ remove it from the collection and label the new refined collection as \mathcal{C}_1 . From this refined collection, if there exists a circle γ such that $|\{\gamma' \mid (\gamma, \gamma') \in \tau(\mathcal{C}_1)\}| < c_1 N^{1/2}$ we remove it and label the remaining collection as \mathcal{C}_2 . After repeating this process M times until there are no more circles γ that satisfy $|\{\gamma' \mid (\gamma, \gamma') \in \tau(\mathcal{C}_M)\}| < c_1 N^{1/2}$, at each step removing a circle that contributes only a small number of tangencies, we attain a collection \mathcal{C}_M . We claim that $|\tau(\mathcal{C}_M)| \gtrsim N^\alpha$, and that $|\mathcal{C}_M| \gtrsim N^{\alpha/2}$.

For the first claim, observe that at each step i we are reducing $\tau(\mathcal{C}_i)$ by at most $c_1 N^{\alpha-1}$. Thus,

$$\begin{aligned} |\tau(\mathcal{C}_M)| &\geq |\tau(\mathcal{C}_0)| - M c_1 N^{\alpha-1} \\ &> c_0 N^\alpha - M c_1 N^{\alpha-1} \\ &> c_0 N^\alpha - c_1 N^\alpha \\ |\tau(\mathcal{C}_M)| &> \frac{c_0}{2} N^\alpha. \quad (\text{by choosing } c_1 = c_0/2) \end{aligned}$$

We now provide a lower bound on the cardinality of our refined set \mathcal{C}_M . We have the trivial inequality $|\tau(\mathcal{C}_M)| \leq |\mathcal{C}_M|^2$. Combining this with the result above, we attain $|\mathcal{C}_M| \gtrsim |N|^{\alpha/2}$. \square

We can now prove the main theorem using a strategy is similar to that of Chapter 2. We assume that there are $\gtrsim N^{3/2}$ points of incidence and argue by contradiction. We will find a polynomial which has a zero set containing many points of incidence between our curves, and then use the structure of the problem to claim that this polynomial's zero set must then also contain all the circles in the collection. We do this to achieve a contradiction as the degree of our polynomial will be too low to contain this many circles in its zero set.

Proof of Theorem 6.0.2. Given an arbitrary collection of circles \mathcal{C} with $\gtrsim N^{3/2}$ tangencies, Lemma 6.2.1 with $\alpha = \frac{3}{2}$ allows us to reduce to a collection $\Gamma \subset \mathcal{C}$ where each circle is tangent to $\gtrsim N^{1/2}$ other circles. After applying a small rotation, we can assume that the tangent line at each point of tangency does not point vertically in the y -direction. Let $\beta(\Gamma) = \{\beta(\gamma) \mid \gamma \in \Gamma\}$, where β is the lifting transform defined earlier. Recall from Lemma 6.1.1 that two circles γ_1 and γ_2 are tangent if and only if $\beta(\gamma_1) \cap \beta(\gamma_2) \neq \emptyset$.

Suppose $(x, y, z) \in \beta(\gamma_1) \cap \beta(\gamma_2)$ for some $\gamma_1 \neq \gamma_2$. Then we claim that

$$(0, 0, 1) \in \text{span}(T_{(x,y,z)}\beta(\gamma_1), T_{(x,y,z)}\beta(\gamma_2)).$$

In other words, at the intersection of $\beta(\gamma_1)$ and $\beta(\gamma_2)$ their tangent vectors span a vertical subspace of \mathbb{R}^3 . We can establish this by examining a parameterisation of γ_1 and γ_2 in the neighbourhood of (x, y) . Define $f_i(t)$, $i \in \{1, 2\}$ such that $(x+t, f_i(t))$ is a parameterisation of γ_i in the neighbourhood of (x, y) for all t in a small neighbourhood of 0. In particular $f_i(0) = y$. Since $(x+t, f_i(t))$ satisfies the equation of a circle γ_i we have the relation

$$(x+t-x_{\gamma_i})^2 + (f_i(t)-y_{\gamma_i})^2 = r_{\gamma_i}^2.$$

Taking the first derivative at $t=0$ and rearranging we attain

$$\frac{df_i}{dt}(0) = -\frac{x-x_{\gamma_i}}{y-y_{\gamma_i}}.$$

Since γ_1 is tangent to γ_2 at (x, y) the slopes of their tangent lines coincide at that point so $\frac{df_1}{dt}(0) = \frac{df_2}{dt}(0)$. Now taking the second derivative at $t=0$ and rearranging

$$\frac{d^2 f_i}{dt^2}(0) = -\frac{(y-y_{\gamma_i})^2 + (x-x_{\gamma_i})^2}{(y-y_{\gamma_i})^3} = -\frac{r_{\gamma_i}^2}{(y-y_{\gamma_i})^3}.$$

Since γ_1 and γ_2 are distinct circles, $\frac{d^2 f_1}{dt^2}(0) \neq \frac{d^2 f_2}{dt^2}(0)$. (In the case $r_{\gamma_1} = r_{\gamma_2}$, notice that the tangency must be external and hence $\text{sgn}(y-y_{\gamma_1})^3 \neq \text{sgn}(y-y_{\gamma_2})^3$)

In the neighbourhood of (x, y, z) , $\beta(\gamma_i)$ is parametrised by $\left(t, f_i(t), \frac{df_i}{dt}(t)\right)$ as the slope of the tangent to the circle is given by $\frac{df_i}{dt}(t)$. It follows that the tangent vector $\left(1, \frac{df_i}{dt}(0), \frac{d^2 f_i}{dt^2}(0)\right)$ generates the vertical space $T_{(x,y,z)}\beta(\gamma_i)$. Thus

$$\begin{aligned} (0, 0, 1) &\in \text{span} \left(\left(1, \frac{df_1}{dt}(0), \frac{d^2 f_1}{dt^2}(0)\right) - \left(1, \frac{df_2}{dt}(0), \frac{d^2 f_2}{dt^2}(0)\right) \right) \\ &\subset \text{span} (T_{(x,y,z)}\beta(\gamma_1), T_{(x,y,z)}\beta(\gamma_2)). \end{aligned} \quad (6.1)$$

We will now interpolate $\sim N^{3/2}$ incidence points with a minimal polynomial whose degree we can control, and show that if this contains too many intersections it must also contain the curves. Then due to the tangent vectors spanning the z -axis at the incidence points we will be able to achieve a contradiction.

For each $\beta(\gamma) \in \beta(\Gamma)$ define a collection K_γ as a set of $\sim N^{1/2}$ incidence points of $\beta(\gamma)$. Then $|\cup_\Gamma K_\gamma| \lesssim N^{3/2}$. Let $P \in \mathbb{R}[x, y, z]$ be a non-zero polynomial of minimal degree that vanishes on all the incidence points in $\cup_\Gamma K_\gamma$. This polynomial interpolates $\lesssim N^{3/2}$ points in \mathbb{R}^3 , so by Lemma 2.2.2 the degree of P is $\lesssim (N^{3/2})^{\frac{1}{3}} = N^{1/2}$.

Due to our refinement each $\gamma \in \Gamma$ is tangent to $\gtrsim N^{1/2}$ circles and each of these tangencies occurs at a distinct point by the non-degeneracy condition. Hence we have that P vanishes at $\gtrsim N^{1/2}$ points of each curve in $\beta(\Gamma)$. By Bézout's Lemma (Lemma 6.1.2) we have that P vanishes on all curves in $\beta(\Gamma)$ as

$$\deg(P) \deg(\beta(\gamma)) \lesssim \#\{P \cap \beta(\gamma)\}$$

for suitable choice of implicit constants.

By Equation 6.1, if (x, y, z) is a point where two curves from $\beta(\Gamma)$ intersect, then $\partial_z P(x, y, z) = 0$. Thus by the same Bézout argument $\partial_z P$ is also a polynomial which vanishes on all curves in $\beta(\Gamma)$.

Since P was a non-zero polynomial of minimal degree that vanishes on all points of incidence in $\beta(\Gamma)$, we must conclude $\partial_z P \equiv 0$ and hence we must have that $P(X, Y, Z) = Q(X, Y)$ for some $Q \in \mathbb{R}[X, Y]$ with $\deg Q = \deg P \lesssim N^{1/2}$. But this implies that there must be at least of the $N^{3/4}$ circles in our original collection which are contained in $Z(Q)$. This is a contradiction, as Q has degree $\sim N^{1/2}$ whereas $\cup \gamma$ has degree $\gtrsim N^{3/4}$. Thus $\beta(\Gamma)$ has $\lesssim N^{3/2}$ curve-curve incidences, so we can conclude that Γ has $\lesssim N^{3/2}$ tangencies. \square

6.3 A New Proof via Polynomial Partitioning for Varieties

In this section we shall present a new proof of Theorem 6.0.2. There are a few key differences between the two methods of proof. In the following proof, it is no longer required to ensure the uniform distribution of tangent points among each circle in our collection, so we will not have to use Lemma 6.2.1. Secondly, in a similar fashion to the proof of the Szemerédi-Trotter Theorem (Theorem 5.0.1) in the previous chapter we will partition \mathbb{R}^3 using a polynomial of controlled degree and leverage our trivial bound (Lemma 6.0.1) in each cell. Slight care will be needed to deal with intersections of curves in the zero set, which we do via a recursive argument.

Definition 6.3.1. An **algebraic variety** in \mathbb{R}^n is defined as $Z(Q_1, \dots, Q_k)$ for some non-zero polynomials Q_1, \dots, Q_k . The **dimension of a variety** $Z(Q_1, \dots, Q_k)$ is the dimension of the tangent spaces at its regular points (these are the points for which the matrix formed by the gradient vectors $\nabla Q_1, \dots, \nabla Q_k$ has maximal rank).

In particular, we refer to algebraic varieties of dimension 1 as **curves**.

We begin by introducing the main tool for this proof, an extension of the polynomial partitioning theorem (Theorem 5.3.3) to algebraic varieties instead of just points (which are 0-dimensional varieties).

Lemma 6.3.1 (Polynomial Partitioning for Algebraic Varieties). *If Γ is a finite set of k -dimensional varieties in \mathbb{R}^n and D any degree, there exists a non-zero polynomial P of degree at most D such that each disjoint open set of $\mathbb{R}^n \setminus Z(P)$ intersects $\lesssim D^{k-n} |\Gamma|$ varieties of Γ .*

We shall not prove this here. A proof of this result can be found as Theorem 0.3 in [Gut15], the original paper of Guth presenting the result. The proof is in the same spirit as the proof Theorem 5.3.3 given in Chapter 5, however some additional difficulties present themselves and a correct proof requires some care. One should notice that the $k = 0$ case above recovers Theorem 5.3.3.

Proof of Theorem 6.0.2 by Polynomial Partitioning. Relabel \mathcal{C} as Γ for convenience. Again, we perform the lifting transform β on each $\gamma \in \Gamma$. We have a collection of N 1-dimensional varieties in \mathbb{R}^3 upon which we use our polynomial partitioning lemma (Lemma 6.3.1) to find a polynomial P such that each cell of $\mathbb{R}^3 \setminus Z(P)$ intersects $\lesssim ND^{-2}$ varieties. $\mathbb{R}^3 \setminus Z(P)$ partitions the space into $\sim D^3$ cells. Let us label the interior of each of these cells as Ω_i for $0 \leq i \lesssim D^3$, and further label the set of varieties in Γ that intersect a given cell Ω_i as Γ_i .

We can now define the following complementary sets based on whether the variety is contained entirely in $Z(P)$:

$$\begin{aligned} C_1 &= \{\beta(\gamma) \mid \beta(\gamma) \not\subset Z(P)\}, \\ C_2 &= \{\beta(\gamma) \mid \beta(\gamma) \subset Z(P)\}. \end{aligned}$$

Notice here that $\beta(\Gamma) = C_1 \cup C_2$. Recalling the correspondence between curve-incidences and circle-tangencies, we define the following incidence sets for any subcollections C and C' of $\beta(\Gamma)$

$$I(C, C') = \{(\beta(\gamma), \beta(\gamma')) \mid \beta(\gamma) \in C', \beta(\gamma') \in C, \beta(\gamma) \cap \beta(\gamma') \neq \emptyset\}.$$

Hence we can express $|\tau(\Gamma)|$ as the sum

$$|\tau(\Gamma)| = |I(C_1, C_1)| + 2|I(C_1, C_2)| + |I(C_2, C_2)|.$$

We now proceed to bound the cardinality for each of these sets. For $I(C_1, C_1)$ will use our trivial bound (Lemma 6.0.1) and Bézout's Lemma (Lemma 6.1.2). Bounding the cardinality for $I(C_1, C_2)$ is again straightforward by Bézout's Lemma. The interesting case here is $I(C_2, C_2)$, where we will be forced to argue via a recursive argument.

We begin with $I(C_1, C_1)$, that is, the intersections that occur between varieties not entirely contained in the zero set. We proceed by partitioning $I(C_1, C_1)$ according to whether the incidence occurs within a cell or on $Z(P)$, explicitly

$$\begin{aligned} I(C_1, C_1) &= \{(\beta(\gamma), \beta(\gamma')) \mid \beta(\gamma), \beta(\gamma') \in C_1, \beta(\gamma) \cap \beta(\gamma') \in \mathbb{R}^3 \setminus Z(P)\} \\ &\quad \cup \{(\beta(\gamma), \beta(\gamma')) \mid \beta(\gamma), \beta(\gamma') \in C_1, \beta(\gamma) \cap \beta(\gamma') \in Z(P)\}. \end{aligned}$$

Hence we have

$$\begin{aligned} |I(C_1, C_1)| &= \sum_{\beta(\gamma), \beta(\gamma') \in C_1} \mathbb{1}[\beta(\gamma) \cap \beta(\gamma') \in \mathbb{R}^3 \setminus Z(P)] + \sum_{\beta(\gamma), \beta(\gamma') \in C_1} \mathbb{1}[\beta(\gamma) \cap \beta(\gamma') \in Z(P)] \\ &= \sum_{\beta(\gamma), \beta(\gamma') \in C_1} \sum_i \mathbb{1}[\beta(\gamma) \cap \beta(\gamma') \in \Omega_i] + \sum_{\beta(\gamma), \beta(\gamma') \in C_1} \mathbb{1}[\beta(\gamma) \cap \beta(\gamma') \in Z(P)] \\ &= \sum_i \sum_{\beta(\gamma), \beta(\gamma') \in \Gamma_i} \mathbb{1}[\beta(\gamma) \cap \beta(\gamma') \in \Omega_i] + \sum_{\beta(\gamma), \beta(\gamma') \in C_1} \mathbb{1}[\beta(\gamma) \cap \beta(\gamma') \in Z(P)]. \end{aligned}$$

Using our trivial bound (Lemma 6.0.1) and the fact that there are $\lesssim ND^{-2}$ varieties intersecting a given cell we attain

$$\begin{aligned} &\lesssim \sum_i \left(\frac{N}{D^2} \right)^2 + \sum_{\beta(\gamma), \beta(\gamma') \in C_1} \mathbb{1}[\beta(\gamma) \cap \beta(\gamma') \in Z(P)] \\ &\sim D^3 N^2 D^{-4} + \sum_{\beta(\gamma) \in C_1} \sum_{\beta(\gamma') \in C_1} \mathbb{1}[\beta(\gamma) \cap \beta(\gamma') \in Z(P)]. \end{aligned}$$

For each fixed $\beta(\gamma) \in C_1$ we see by Bézout's Lemma (Lemma 6.1.2) that $\beta(\gamma)$ can intersect $Z(P)$ in at most $\lesssim D$ points. Since all the intersections in the latter sum occur in $Z(P)$, we have that each $\beta(\gamma)$ contributes at most $\lesssim D$ to the sum, and summing over all possible $\beta(\gamma)'s \in C_1$ we obtain

$$\begin{aligned} &\lesssim N^2 D^{-1} + \sum_{\beta(\gamma) \in C_1} D \\ &\lesssim N^2 D^{-1} + ND. \end{aligned}$$

The argument for $I(C_1, C_2)$ is similar to the one above, however there are now no intersections happening inside the cells by the definition of C_2 , hence we have

$$|I(C_1, C_2)| = \sum_{\beta(\gamma) \in C_1} \sum_{\beta(\gamma') \in C_2} \mathbb{1}[\beta(\gamma) \cap \beta(\gamma') \in Z(P)].$$

Again thanks to Bézout's Lemma (Lemma 6.1.2) we have that the latter sum is $\lesssim D$. Hence

$$|I(C_1, C_2)| \lesssim \sum_{\beta(\gamma) \in C_1} D \leq ND.$$

Finally we need to handle the incidences between curves contained entirely in the zero-set, $I(C_2, C_2)$. This is a slightly more technical undertaking than the previous two cases. We proceed by considering the class of polynomials whose zero set contains every curve in C_2 , formally defined as

$$\mathcal{P}_0 = \{R(X, Y, Z) \in \mathbb{R}[X, Y, Z] \mid \forall \beta(\gamma) \in C_2, \beta(\gamma) \subset Z(R)\}.$$

\mathcal{P}_0 is non-empty, as by the definition of C_2 it at least contains the partitioning polynomial P . Let us take the polynomial in \mathcal{P}_0 of minimal degree and label it P_0 . As seen in the proof in Section 6.2, since the curves are entirely contained in $Z(P_0)$, and at a point of incidence we have $\partial_z P_0 = 0$, we conclude $Z(\partial_z P_0)$ contains all the points in the set $I(C_2, C_2)$. Either $\partial_z P_0 \equiv 0$ or it is not.

In the first case, we must have that $P_0(X, Y, Z) = Q(X, Y)$ for some polynomial $Q \in \mathbb{R}[X, Y]$. $Z(Q)$ contains all the circles γ such that $\beta(\gamma) \in C_2$ (to see this, notice that

$\beta(\gamma)$ restricted to \mathbb{R}^2 is just γ). By Bézout's Lemma, $Z(Q)$ contains at most $\lesssim \deg \partial_z P_0$ circles, hence we can trivially bound the incidences here by $|I(C_2, C_2)| \leq (\deg \partial_z P_0)^2 \lesssim D^2$. In this case there is nothing left to do as we have bounded $|I(C_2, C_2)|$, so the recursion stops.

In the second case, we must have that $\deg \partial_z P_0 < \deg P_0 \leq D$. Since P_0 was of minimal degree in \mathcal{P}_0 , $Z(\partial_z P_0)$ cannot contain all the curves of C_2 . Therefore we can split the collection into a complementary pair of sets

$$\begin{aligned} C_1^{(1)} &= \{\beta(\gamma) \in C_2 \mid \beta(\gamma) \not\subset Z(\partial_z P_0)\}, \\ C_2^{(1)} &= \{\beta(\gamma) \in C_2 \mid \beta(\gamma) \subset Z(\partial_z P_0)\}. \end{aligned}$$

Using these sets we can express

$$|I(C_2, C_2)| = |I(C_1^{(1)}, C_1^{(1)})| + 2|I(C_1^{(1)}, C_2^{(1)})| + |I(C_2^{(1)}, C_2^{(1)})|$$

Notice that since all points of incidence belong to $Z(\partial_z P_0)$, the set $I(C_1^{(1)}, C_1^{(1)})$ is empty. Hence

$$= |I(C_1^{(1)}, C_2)| + |I(C_2^{(1)}, C_2^{(1)})|.$$

It follows from Bézout's Lemma via the same argument given before that $|I(C_1^{(1)}, C_2)| \lesssim |C_1^{(1)}| D$, and we will now repeat this process to bound $|I(C_2^{(1)}, C_2^{(1)})|$.

We now move into the next step of the recursion. Now we consider the class of polynomials

$$\mathcal{P}_1 = \{R(X, Y, Z) \in \mathbb{R}[X, Y, Z] \mid \forall \beta(\gamma) \in C_2^{(1)}, \beta(\gamma) \subset Z(R)\}.$$

\mathcal{P}_1 is non-empty, as by the definition of $C_2^{(1)}$ it at least contains the polynomial $\partial_z P_0$ (and hence the minimal degree polynomial in \mathcal{P}_1 is of degree at most $D - 1$). Let us take the polynomial in \mathcal{P}_1 of minimal degree and label it P_1 . Again we have that $Z(\partial_z P_1)$ contains all the points of $I(C_2^{(1)}, C_2^{(1)})$. Again by the law of the excluded middle we have either $\partial_z P_1 \equiv 0$ or it is not.

In the first case, again we must have that $P_1(X, Y, Z) = Q(X, Y)$ for some polynomial $Q \in \mathbb{R}[X, Y]$. $Z(Q)$ contains all the circles γ such that $\beta(\gamma) \in C_2$, of which there are $\lesssim \deg \partial_z P_0$. Hence we can trivially bound the incidences here by

$$|I(C_2^{(1)}, C_2^{(1)})| \leq (\deg \partial_z P_1)^2 \lesssim D^2.$$

If this is the case the recursion ends.

In the second case, we must have that $\deg \partial_z P_1 < \deg P_1 < \deg \partial_z P_0 \lesssim D$. Again, since P_1 was of minimal degree in \mathcal{P}_1 , $Z(\partial_z P_1)$ cannot contain all the curves of C_2 . Hence

we now define a new complementary pair of sets as

$$\begin{aligned} C_1^{(2)} &= \{\beta(\gamma) \in C_2^{(1)} \mid \beta(\gamma) \not\subset Z(\partial_z P_1)\}, \\ C_2^{(2)} &= \{\beta(\gamma) \in C_2^{(1)} \mid \beta(\gamma) \subset Z(\partial_z P_1)\}. \end{aligned}$$

Using these sets we can express

$$\begin{aligned} \left| I(C_2^{(1)}, C_2^{(1)}) \right| &= \left| I(C_1^{(2)}, C_1^{(2)}) \right| + \left| I(C_1^{(2)}, C_2^{(2)}) \right| + \left| I(C_2^{(2)}, C_2^{(2)}) \right| \\ &= \left| I(C_1^{(2)}, C_2^{(1)}) \right| + \left| I(C_2^{(2)}, C_2^{(2)}) \right|. \end{aligned}$$

We have that $\left| I(C_1^{(2)}, C_2^{(1)}) \right| \lesssim \left| C_1^{(2)} \right| D$. Again, we can continue the recursion to bound $\left| I(C_2^{(2)}, C_2^{(2)}) \right|$.

In the i -th step of our recursion, we define

$$\mathcal{P}_i = \{R(X, Y, Z) \in \mathbb{R}[X, Y, Z] \mid \forall \beta(\gamma) \in C_2^{(i)}, \beta(\gamma) \subset Z(R)\}.$$

This will always be non-empty as by the definition of $C_2^{(i)}$, it always contains at least the polynomial $\partial_z P_{i-1}$. The polynomial of minimal degree in \mathcal{P}_i , denoted P_i , is of strictly smaller degree than P_{i-1} . Hence our recursion must end after at most D steps. At each step, either $\partial_z P_i \equiv 0$ or it is not.

If $\partial_z P_i \equiv 0$, we can bound the contribution of this step by

$$\left| I(C_2^{(i)}, C_2^{(i)}) \right| \lesssim D^2$$

and our recursion ends.

In the second case, we again define the complementary sets

$$\begin{aligned} C_1^{(i+1)} &= \{\beta(\gamma) \in C_2^{(i)} \mid \beta(\gamma) \not\subset Z(\partial_z P_i)\}, \\ C_2^{(i+1)} &= \{\beta(\gamma) \in C_2^{(i)} \mid \beta(\gamma) \subset Z(\partial_z P_i)\}. \end{aligned}$$

As before, these will lead to a contribution of $\lesssim \left| C_1^{(i+1)} \right| D + \left| I(C_2^{(i+1)}, C_2^{(i+1)}) \right|$. The process continues on the latter term.

Let us collect the contributions from each step of our recursion. At the terminal step of the process, say step J , we collect a term of $\lesssim D^2$. Collecting the terms preceeding this, we are able to bound the contribution by

$$\sum_{i=0}^{J-1} \left| C_1^{(i+1)} \right| D \lesssim ND$$

by using the fact that all $C_1^{(i+1)}$ are disjoint subsets of C_2 and that C_2 contains at most

N curves, we can conclude that

$$|I(C_2, C_2)| \lesssim D^2 + ND.$$

Adding together our bounds we achieve

$$\begin{aligned} |I(C_1, C_1)| + 2|I(C_1, C_2)| + |I(C_2, C_2)| &\lesssim N^2 D^{-1} + ND + ND + ND + D^2 \\ &\sim N^2 D^{-1} + ND + D^2. \end{aligned}$$

We now optimise D by setting $N^2 D^{-1} \sim ND$ which gives $D \sim N^{1/2}$ and hence our bounds become

$$|I(C_1, C_1)| + |I(C_1, C_2)| + |I(C_2, C_2)| \lesssim N^2 N^{-1/2} + NN^{1/2} + (N^{1/2})^2 \sim N^{3/2}.$$

□

Remark. *Since polynomial partitioning exploits the topology of \mathbb{R}^n in an essential manner, it would have been reasonable to expect an improvement over the exponent $3/2$ via the argument above. Alas, this was not the case.*

6.4 Considerations for the Case of Sphere Tangencies in \mathbb{R}^3

Considering the analogous problem in \mathbb{R}^3 , two new difficulties emerge. Firstly, our current non-degeneracy condition is not sufficient as we have new degenerate cases to consider where the trivial bound of $O(N^2)$ is achieved.

Example 6.3 (Degenerate case in \mathbb{R}^3 , [Zah22]). *Consider the two collections of N spheres illustrated in Figure 6.2. Under the union of these collections, each sphere in the first collection is tangent to $\sim N$ spheres in the second collection at N distinct points. Hence the union is a collection of $2N$ spheres that produces N^2 tangencies. Note that for each sphere in the first collection, all of its tangencies with the second collection occur along a circle on the surface of the sphere.*

Due to cases like these a new condition is required. A natural choice given this example would seem to be restricting our collections to those which do not contain too many tangencies along any one low-degree algebraic curve on the surface of a sphere.

Secondly, our simple Bézout lemma is no longer enough. An analogous lifting of a sphere $\gamma \subset \mathbb{R}^3$ is given by




Figure 6.2: Two collections of N spheres in \mathbb{R}^3 , whose union satisfies the non-degeneracy condition the earlier section yet achieves N^2 tangencies.

$$\beta(\gamma) = \left\{ (x, y, z, v, w) \in \mathbb{R}^5 \mid (x - x_\gamma)^2 + (y - y_\gamma)^2 + (z - z_\gamma)^2 = r_\gamma^2, \right. \\ \left. v = -\frac{x - x_\gamma}{z - z_\gamma}, w = -\frac{y - y_\gamma}{z - z_\gamma} \right\}.$$

Notice that $\beta(\gamma)$ is a two-dimensional algebraic surface in \mathbb{R}^5 . An interpolating polynomial's zero set $Z(P)$ such as the one used in Ellenberg-Solymosi-Zahl's proof would be a 4-dimensional surface in \mathbb{R}^5 . These two surfaces could intersect along a 1-dimensional variety (i.e. at infinitely many points), hence we can no longer argue that $\beta(\gamma) \subset Z(P)$ due to $Z(P)$ containing too many points of $\beta(\gamma)$.

The best example of a lower bound for such tangencies is produced in a similar fashion to the circles case (Example 6.2). In place of the Szemerédi-Trotter theorem we appeal to a result of [EGS90], later refined by the authors in [AS07] to remove a subpolynomial term.¹

Theorem 6.4.1 ([EGS90], [AS07]). *Let \mathcal{S} be a finite set of points in \mathbb{R}^3 and Π a finite set of collinear planes in \mathbb{R}^3 such that no three are collinear. Then the incidence set $I(\mathcal{S}, \Pi)$ obeys*

$$|I(\mathcal{S}, \Pi)| \lesssim |\mathcal{S}|^{4/5} |\Pi|^{3/5} + |\mathcal{S}| + |\Pi|.$$

The above bound is indeed tight as shown by a construction in [BK03]. We now construct our lower bound utilising geometric inversion in \mathbb{R}^3 analogously to Example 6.2.

Example 6.4. *Let \mathcal{S}, Π be collections of N points and N non-collinear planes respectively such that the Theorem 6.4.1 is asymptotically tight. In particular, they determine $\sim N^{7/5}$*

¹The original bound in [EGS90] was multiplied by a factor of $|\mathcal{S}|^\delta |\Pi|^\delta$ for some $\delta > 0$.

point-plane incidences. Let C_1 denote the collection of unit spheres centred at the points of \mathcal{S} . Let C_2 denote the collection of planes obtained by translating each plane $\pi \in \Pi$ one unit in the π^\perp direction. If $(p, \pi) \in \mathcal{S} \times \Pi$ is a point-plane incidence from our original collection, then the corresponding sphere-plane pair will be tangent. Performing an inversion transform about a point that does not lie in any of the circles or planes, our collection $\mathcal{C} = C_1 \cup C_2$ becomes a collection of $2N$ spheres that determine $N^{3/2}$ tangencies.

In [ESZ16], the authors conjectured a bound for higher dimensional spheres:

Conjecture 4. Let \mathcal{C} be a collection of $(d-1)$ -spheres in \mathbb{R}^d . Then the number of distinct points of tangencies is $\lesssim |\mathcal{C}|^{\frac{2d-1}{d}}$.

This is a somewhat weak conjecture, as the evidence to form it is simply the extension of the β -lifting technique into higher dimensions. The β -lifting of spheres in \mathbb{R}^d produces curves in \mathbb{R}^{2d-1} , and it is the ratio of these dimensions that authors claim to be a bound for the exponent. As such, there is no evidence to suggest that it should be tight.

References

- [Alo99] Noga Alon. Combinatorial Nullstellensatz. volume 8, pages 7–29. 1999. Recent trends in combinatorics (Mátraháza, 1995).
- [Alo03] Noga Alon. Non-constructive proofs in Combinatorics. Online Lecture Notes (<http://www.math.tau.ac.il/~nogaa/PDFS/nocon.pdf>), 2003.
- [AS07] Roel Apfelbaum and Micha Sharir. Large complete bipartite subgraphs in incidence graphs of points and hyperplanes. *SIAM J. Discrete Math.*, 21(3):707–725, 2007.
- [BK03] Peter Brass and Christian Knauer. On counting point-hyperplane incidences. volume 25, pages 13–20. 2003. Special issue on the European Workshop on Computational Geometry—CG01 (Berlin).
- [Bou91] Jean Bourgain. Besicovitch type maximal operators and applications to Fourier analysis. *Geom. Funct. Anal.*, 1(2):147–187, 1991.
- [CEG⁺90] Bernard Chazelle, Herbert Edelsbrunner, Leonidas J. Guibas, et al. Counting and cutting cycles of lines and rods in space. In *31st Annual Symposium on Foundations of Computer Science, Vol. I, II (St. Louis, MO, 1990)*, pages 242–251. IEEE Comput. Soc. Press, Los Alamitos, CA, 1990.
- [Dav71] Roy O. Davies. Some remarks on the Kakeya problem. *Proc. Cambridge Philos. Soc.*, 69:417–421, 1971.
- [Dvi09] Zeev Dvir. On the size of Kakeya sets in finite fields. *J. Amer. Math. Soc.*, 22(4):1093–1097, 2009.
- [EGS90] Herbert Edelsbrunner, Leonidas Guibas, and Micha Sharir. The complexity of many cells in arrangements of planes and related problems. *Discrete Comput. Geom.*, 5(2):197–216, 1990.
- [ESZ16] Jordan S. Ellenberg, Jozsef Solymosi, and Joshua Zahl. New bounds on curve tangencies and orthogonalities. *Discrete Anal.*, pages Paper No. 18, 22, 2016.
- [FS05] Sharona Feldman and Micha Sharir. An improved bound for joints in arrangements of lines in space. *Discrete Comput. Geom.*, 33(2):307–320, 2005.

- [FT81] Robert M. Freund and Michael J. Todd. A constructive proof of Tucker’s combinatorial lemma. *J. Combin. Theory Ser. A*, 30(3):321–325, 1981.
- [Ful98] William Fulton. *Intersection theory*, volume 2 of *Ergebnisse der Mathematik und ihrer Grenzgebiete. 3. Folge. A Series of Modern Surveys in Mathematics [Results in Mathematics and Related Areas. 3rd Series. A Series of Modern Surveys in Mathematics]*. Springer-Verlag, Berlin, second edition, 1998.
- [GK10] Larry Guth and Nets Hawk Katz. Algebraic methods in discrete analogs of the Kakeya problem. *Adv. Math.*, 225(5):2828–2839, 2010.
- [Gut15] Larry Guth. Polynomial partitioning for a set of varieties. *Math. Proc. Cambridge Philos. Soc.*, 159(3):459–469, 2015.
- [Gut16] Larry Guth. *Polynomial methods in combinatorics*, volume 64 of *University Lecture Series*. American Mathematical Society, Providence, RI, 2016.
- [KSS10] Haim Kaplan, Micha Sharir, and Eugenio Shustin. On lines and joints. *Discrete Comput. Geom.*, 44(4):838–843, 2010.
- [Mat03] Jiří Matoušek. *Using the Borsuk-Ulam theorem*. Universitext. Springer-Verlag, Berlin, 2003. Lectures on topological methods in combinatorics and geometry, Written in cooperation with Anders Björner and Günter M. Ziegler.
- [Mic10] Mateusz Michałek. A short proof of combinatorial Nullstellensatz. *Amer. Math. Monthly*, 117(9):821–823, 2010.
- [Pal20] Julius Pal. Über ein elementares variationsproblem, det kgl. danske videnskab. selskab. *Math. Fys. Meddel*, 3(2):1–35, 1920.
- [Qui10] René Quilodrán. The joints problem in \mathbb{R}^n . *SIAM J. Discrete Math.*, 23(4):2211–2213, 2009/10.
- [ST83] Endre Szemerédi and William T. Trotter, Jr. Extremal problems in discrete geometry. *Combinatorica*, 3(3-4):381–392, 1983.
- [Szé97] László A. Székely. Crossing numbers and hard Erdős problems in discrete geometry. *Combin. Probab. Comput.*, 6(3):353–358, 1997.
- [Wol95] Thomas Wolff. An improved bound for Kakeya type maximal functions. *Rev. Mat. Iberoamericana*, 11(3):651–674, 1995.
- [Wol99] Thomas Wolff. Recent work connected with the Kakeya problem. In *Prospects in mathematics (Princeton, NJ, 1996)*, pages 129–162. Amer. Math. Soc., Providence, RI, 1999.
- [Zah22] Joshua Zahl. Sphere tangencies, line incidences and Lie’s line-sphere correspondence. *Math. Proc. Cambridge Philos. Soc.*, 172(2):401–421, 2022.

Appendix A

Proof of the Borsuk-Ulam Theorem

Theorem A.0.1 (Borsuk-Ulam (Theorem 5.3.1)). *A map ϕ is said to be antipodal if it obeys $\phi(-x) = -\phi(x)$ for all x in its domain. Suppose $\phi : \mathbb{S}^N \rightarrow \mathbb{R}^N$ is a continuous antipodal mapping. Then the image of ϕ contains 0.*

We present here a combinatorial proof due to Matousek (See [Mat03]).

Let $\|x\|_1$ be the L^1 norm of x . Define B^n as the unit ball with respect to the L_1 norm, that is: $B^n = \{x \in \mathbb{R}^n \mid \|x\|_1 \leq 1\}$.

Definition A.0.1. A simplex is the convex hull of an affinely independent set of vectors in \mathbb{R}^n . The convex hull of an arbitrary subset of vertices of a simplex is called a face of the simplex.

A family of simplexes $\Delta = \{\sigma_1, \sigma_2, \dots, \sigma_m\}$ is called a **simplicial complex** if the following conditions hold:

1. Each non-empty face of any simplex $\sigma \in \Delta$ is also a simplex of Δ .
2. $\sigma_1, \sigma_2 \in \Delta \implies \sigma_1 \cap \sigma_2$ is a face of both σ_1 and σ_2 .

The union of all simplexes in a simplicial complex Δ is denoted by $\|\Delta\|$.

A simplicial complex T is a **special triangulation** of B^n if all the following hold:

1. $\|T\| = B^n$.
2. T is a refinement of the triangulation of B^n given by cutting the coordinate hyperplanes. (In other words, no simplex of T spans over a boundary of an orthant)
3. T is symmetrical around the origin.

To prove the Borsuk-Ulam theorem, we first prove the Tucker Lemma. This lemma can be thought of as the combinatorial analogue of the Borsuk-Ulam Theorem. We present here a constructive proof first published by Freund.[FT81]

Lemma A.0.2 (Tucker Lemma). *Let the vertices of an arbitrary special triangulation T be denoted by labels $\text{lab}(u) \in \{\pm 1, \pm 2, \dots, \pm n\}$ in such a way that the vertices $u \in \partial B^n$ on the boundary satisfies $\text{lab}(-u) = -\text{lab}(u)$. Then there exists a 1-simplex (an edge) in T which is complementary, that is its two vertices x, x' satisfy $\text{lab}(-x) = -\text{lab}(x')$.*

Proof. Let T be a special triangulation of B^n . For a simplex $\sigma \in T$ we set $\text{sgn } \sigma = (\text{sgn } x_1, \text{sgn } x_2, \dots, \text{sgn } x_n)$, where x is an arbitrary point of the interior of σ . This definition is well-posed, as special triangulations refine the orthants of \mathbb{R}^n and thus the signs of each coordinate do not change in the interior of σ . We say σ is **completely labelled** if the following holds for each $0 \leq i \leq n$: if $\text{sgn } (\sigma)_i = 1$, then at least one of the vertices of σ is labelled by the number i , and if $\text{sgn } (\sigma)_i = -1$, then at least one of the vertices of σ is labelled by the number $-i$.

We now define a graph G whose vertices consist of all the completely labelled simplexes of T , and in which two vertices $\sigma, \sigma' \in T$ are connected by an edge if:

- (a) $\sigma, \sigma' \subset \partial B^n = S^{n-1}$ and $\sigma = -\sigma'$, or
- (b) σ is a k -simplex and σ' is one of its $(k-1)$ -face whose vertices are already labelled by all numbers required for a complete labelling of σ .

The degree of a completely labelled simplex is the number of completely labelled simplexes adjacent to it in G .

The simplex $\{0\}$ has degree 1 in G , since it is connected to exactly the edge of the triangulation which is completely labelled by $\text{lab}(0)$. We now prove that any other vertex σ of G has degree 2 except when σ contains a complementary edge. Since a graph cannot contain only one vertex of odd degree, the lemma will be established.

Let $\text{sgn } \sigma$ have k non-zero components, then the dimension of σ is either k or $k-1$.

Suppose first that σ is a $(k-1)$ -simplex. If σ does not lie in S^{n-1} , it is the face of precisely two k -simplices, both completely labelled since σ is. If σ lies in S^{n-1} , it is the face of one completely labelled k -simplex, and it has the other neighbour $-\sigma$ according to condition (a).

If σ is a k -simplex, it has k obligatory labels and one extra label. This label either:

- repeats one of the k obligatory labels in which case σ is adjacent to two of its faces, or
- it is opposite to one of the obligatory labels in which case we have a complimentary edge, or
- it is yet another number not in the k obligatory labels, in which case the neighbours of σ are its completely labelled face and one adjacent simplex of larger dimension determined by the extra label.

In both cases without a complementary edge we have two neighbours. □

Proof of Borsuk-Ulam from Tucker Lemma. Let $f : \mathbb{S}^N \rightarrow \mathbb{R}^N$ be a continuous mapping, and let B^N be the unit ball in an “equator” hyperplane of \mathbb{S}^N . We define $g : B^N \rightarrow \mathbb{R}^N$ by setting $g(x) = f(y) - f(-y)$, where y is the point of the upper hemisphere of \mathbb{S}^N whose vertical projection on B^N is x . The map g is obviously antipodal on $\partial B^N = \mathbb{S}^{N-1}$. For contradiction let us assume that $g(x) \neq 0$ everywhere. From the compactness of the ball we have the existence of an $\varepsilon > 0$ such that $\|g(x)\|_1 \geq \varepsilon$ for all x . Further, a continuous function on a compact set is uniformly continuous, and thus there exists a $\delta > 0$ such that if $\|x - x'\|_1 \leq \delta$ then $\|g(x) - g(x')\|_1 < \varepsilon/N$.

Let us choose a special triangulation T such that the diameter of each of its simplexes is at most δ . We define a labelling of the vertices of T as follows: $|\text{lab}(x)| = i$ if $|g_i(x)| = \max\{|g_1(x)|, \dots, |g_N(x)|\}$ and $\text{sgn lab}(x) = \text{sgn } g_i(x)$ (if the maximum is attained for more than one index, we take the first such index). From the Tucker Lemma we know there exists a complementary edge xx' . Let $\text{lab}(x) = -\text{lab}(x') = i$, then $g(x)_i \geq \varepsilon/N$ and $g(x')_i \leq -\varepsilon/N$. Hence $\|g(x) - g(x')\|_1 \geq 2\varepsilon/N$, a contradiction. Therefore there exists a zero x of the function g . \square